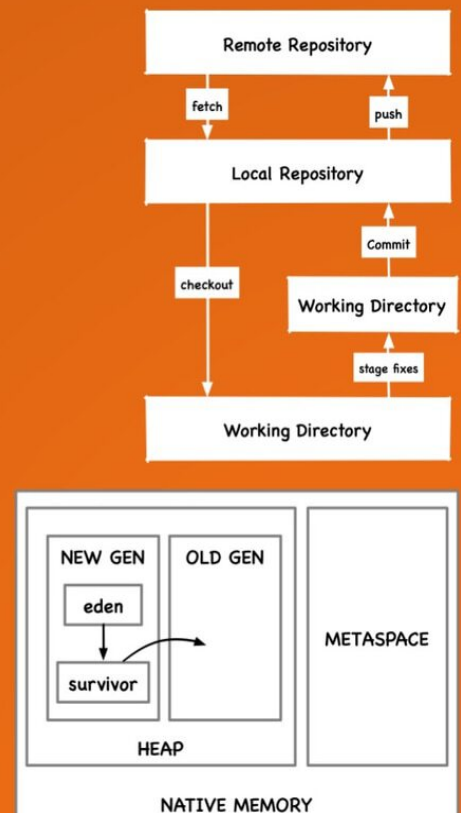



PREPARE FOR YOUR JAVA INTERVIEW IN A DAY!!!

JAVA

INTERVIEW

COURSE



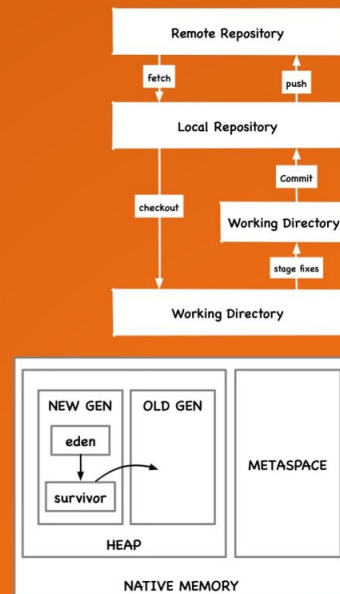

JOLLYM

PREPARE FOR YOUR JAVA INTERVIEW IN A DAY!!!

JAVA

INTERVIEW

COURSE



JOLLYM

JAVA INTERVIEW COURSE

Jolly M

Copyright © 2020 Jolly M

All rights reserved

The characters and events portrayed in this book are fictitious. Any similarity to real persons, living or dead, is coincidental and not intended by the author.

No part of this book may be reproduced, or stored in a retrieval system, or transmitted in any form or by any means, electronic, mechanical, photocopying, recording, or otherwise, without express written permission of the publisher.

ISBN-13: 9781234567890

ISBN-10: 1477123456

Cover design by: Art Painter

Library of Congress Control Number: 2018675309

Printed in the United States of America

CONTENTS

[Title Page](#)

[Copyright](#)

[Introduction](#)

[JAVA FUNDAMENTALS](#)

[Java Program Anatomy.](#)

[Compiling and Executing Java Code in JVM](#)

[Data Types](#)

[Naming Convention](#)

[Object class](#)

[Access Modifiers](#)

[static](#)

[final](#)

[static Initialization Block](#)

[finally.](#)

[finalize\(\)](#)

[Widening vs Narrowing Conversions](#)

[getters and setters](#)

[varargs vs object array.](#)

[Default Interface Method](#)

[Static Interface Method](#)

[Annotations](#)

[Preferences](#)

[Pass by value or Pass by Reference](#)

[OBJECT ORIENTED PROGRAMMING](#)

[Polymorphism](#)

[Overriding](#)

[Overloading](#)

[Abstraction](#)

[Inheritance](#)

[Composition](#)

[FUNDAMENTAL DESIGN CONCEPTS](#)

[Dependency Injection vs Inversion of Control](#)

[Service Locator](#)

[Diamond Problem](#)

[Programming to interface](#)

[Abstract Class vs Interface](#)

[Internationalization and Localization](#)

[Immutable Objects](#)

[Cloning](#)

[DATA TYPES](#)

[NaN](#)

[EnumSet](#)

[Comparing the Types](#)

[Enum vs Public Static Field](#)

[Wrapper Classes](#)

[Auto boxing and Auto unboxing](#)

[BigInteger and BigDecimal](#)

[STRINGS](#)

[String Immutability.](#)

[String Literal vs Object](#)

[String Interning](#)

[String Pool Memory Management](#)

[Immutability - Security Issue](#)

[Circumvent String Immutability](#)

[StringBuilder vs StringBuffer](#)

[Unicode](#)

[INNER CLASSES](#)

[Inner Classes](#)

[Static Member Nested Class](#)

[Local Inner Class](#)

[Non-Static Nested Class](#)

[Anonymous Inner Class](#)

[FUNCTIONAL PROGRAMMING](#)

[Functional Interface](#)

[Lambda Expression](#)

[Pure Functions](#)

[Fluent Interface](#)

[GENERICS](#)

[Generics](#)

[Type Wildcards](#)

[Generic Method](#)

[Java Generics vs Java Array](#)

[Java Array](#)

[Type Erasure](#)

[Co-variance](#)

[Contra-variance](#)

[Co-variance vs Contra-variance](#)

[COLLECTIONS](#)

[Collections](#)

[Collection Fundamentals](#)

[Collection Interfaces](#)

[Collection Types](#)

[Algorithms](#)

[Comparable vs Comparator](#)

[hashCode\(\) and equals\(\)](#)

[HashTable vs HashMap](#)

[Synchronized vs Concurrent Collections](#)

[Iterating Over Collections](#)

[Fail Fast](#)

[ERROR AND EXCEPTION HANDLING](#)

[Exception](#)

[Checked vs Unchecked vs Error](#)

[Exception Handling Best Practices](#)

[try-with-resource](#)

[THREADING](#)

[Threading Terms](#)

[Thread Lifecycle](#)

[Thread Termination](#)

[Implementing Runnable vs Extending Thread](#)

[Runnable vs Callable Interface](#)

[Daemon Thread](#)

[Race Condition and Immutable Object](#)

[Thread Pool](#)

[SYNCHRONIZATION](#)

[Concurrent vs Parallel vs Asynchronous](#)

[Thread Synchronization](#)

[Synchronized Method vs Synchronized Block](#)

[Conditional Synchronization](#)

[Volatile](#)

[static vs volatile vs synchronized](#)

[ThreadLocal Storage](#)

[wait\(\) vs sleep\(\)](#)

[Joining Threads](#)

[Atomic Classes](#)

[Lock](#)

[Synchronisers](#)

[Executor Framework](#)

[Fork-Join](#)

[REFLECTION](#)

[Reflection](#)

[Drawbacks of reflection](#)

[DATA INTERCHANGE](#)

[JSON](#)

[MEMORY MANAGEMENT](#)

[Stack vs Heap](#)

[Heap Fragmentation](#)

[Object Serialization](#)

[Garbage Collection](#)

[Memory Management](#)

[Weak vs Soft vs Phantom Reference](#)

[UNIT TESTING](#)

[Why Unit Testing?](#)

[Unit vs Integration vs Regression vs Validation Testing](#)

[Testing Private Members](#)

[Mocking](#)

[JAVA DEVELOPMENT TOOLS](#)

[Git](#)

[Maven](#)

[Jenkins](#)

[JAVA INTERVIEW COURSE](#)

INTRODUCTION

Java Interview Course book is written based on the experience of many professional Software Architects who have conducted hundreds of technical interviews for their organization.

Java Interview Course covers most of the Java concepts that applicants are expected to have hands-on experience and must have thorough conceptual understating. This book also includes lots of snippets of code and figures to explain the concepts.

A clear understanding of the principles of software design and language is important as the same questions can be discussed during an interview in several different ways.

Interviewers are interested not only in hearing the correct response but also in your opinion and thoughts on the issue. Your approach towards interpreting the questions and articulating your thoughts plays a crucial role in your success.

To build the courage to face such an interview, you need to spend a decent deal of time coding, reading technical topics, and talking to peers about it.

I hope this book helps you plan not just for your next interview but also for your daily task of developing software.

All the best!!!

JAVA FUNDAMENTALS

JAVA PROGRAM ANATOMY

The code snippet below depicts the anatomy of a simple Java program.

```
package com.jtech.basics;

import static java.lang.System.out;

public class HelloWorld {

    static String greeting = "Hello World!";

    public static void main(String [] args){
        out.println(greeting);
    }
}
```

The diagram illustrates the components of the Java code snippet with red arrows pointing to specific parts:

- package**: Points to `com.jtech.basics;`
- import**: Points to `import static java.lang.System.out;`
- type**: Points to `public class HelloWorld`
- field**: Points to `static String greeting = "Hello World!";`
- modifier**: Points to `public static void`
- method**: Points to `main`
- argument**: Points to `String []`
- parameter**: Points to `args`

- **Package** - Represents logical grouping of similar types into namespaces. It prevents naming collisions and access protection.
- **Import** - Imports the package, so that classes can be used in the code by their unqualified names.
- **Class** - Represents a type template having properties and methods.
- **Field** - Represents a member used for holding values.
- **Method** - Represents an operation/behavior of the class.


- **Modifier** - Specifies the access control level of a class and or its members.
- **Parameter** - Specifies the variable declared in the method definition.
- **Argument** - Specifies the data passed to the method parameters.

Questions

- What is a package?
- Why do you need to import packages?
- Why do you need to specify access modifiers?
- What is a static import?
 - *static import enables access to static members of a class without the need to qualify it by the class name.*
- What is the difference between argument and parameter?

COMPILING AND EXECUTING JAVA CODE IN JVM

Java program compilation and execution steps



1. Java compiler compiles the Java source code (.java file) into a binary format known as **bytecode** (.class file). Java bytecode is a platform-independent instruction set, which contains instructions (opcode) and parameter information.

2. **Bytecode** is translated by the operating system specific **Java Virtual Machine (JVM)** into the platform-specific machine code.
3. **The Class loader** in JVM loads the binary representation of the classes into memory.
4. **The Execution engine** in JVM executes the byte code and generates operating system specific machine instructions. These machine code instructions are executed directly by the central processing unit (CPU).

Questions

- Explain the process of Java code compilation and execution?
- What is bytecode?
- What is the difference between bytecode and source code?
- What is the machine code?
- What is the difference between bytecode and machine code?
- What is JVM? Is it different or the same for different operating systems?
- What are the major components of JVM?
- What is the role of the class loader in JVM?
- What is the role of the execution engine in JVM?
- What are machine instructions?

DATA TYPES

Primitive Types

- Primitive data types are `byte` , `boolean` , `char` , `short` , `int` , `float` , `long` and `double` .
- Primitive types always have value; if not assigned, it will have a default value.
- A `long` value is suffixed with `L` (or `l`) to differentiate it from `int` .
- A `float` value is suffixed with `F` (or `f`) to differentiate it from `double` . Similarly `double` is suffixed with `D` (or `d`)
- A `char` is unsigned and represent `Unicode` values.
- When a primitive type is assigned to another variable, a copy is created.

Reference Types

- All non-primitive types are reference types.
- Reference types are also usually known as `objects` . Though reference types refer to an object in memory.
- An unassigned object of reference type will have `null` as the default value.
- Objects have variables and methods, which define their state and behavior.
- When a reference is assigned to another reference, both points to the same object.

Questions

- What are primitive data types?

- If a variable of primitive data type is not assigned, what does it contain?
- Why do we suffix L with long, F with Float and D with a double?
- What happens when you assign a variable of a primitive data type to another variable of the same type?
- What are the reference data types?
- What happens when you assign a variable of the reference data type to another variable of the same reference type?
- What are the differences between the primitive data types and the reference data types?
- What is the purpose of variables and methods in a reference type?
- If a variable of reference data type is not assigned, what does it contain?

NAMING CONVENTION

Camel Case vs Pascal Case

Camel Case is the practice of writing composite words such that the first letter in each word is capitalized, like `BorderLength`; it is also known as **Pascal Case** or **upper Camel Case** . But in the programming world, Camel case generally starts with the lower case letter, like `borderLength`; it is also known as **lower Camel Case** . For this discussion, let's consider the format `BorderLength` as **Pascal Case** and the format `borderLength` as **Camel Case** .

The **Naming convention** is a set of rules that govern the naming for the identifiers, which represents interface, class, method, variable, and other entities. Although the choice and implementation of the naming conventions often become a matter of debate.

Standard naming convention improves the code readability, review, and overall understanding.

Interface

- The name should be **Pascal Case** .
- The name should be an adjective if it defines behavior, otherwise noun.

```
public interface Runnable
```

Class

- The name should be **Pascal Case** .
- The name should be a noun, as a **class** represents some real-world object.

```
public class ArrayList
```

Method

- The name should be Camel Case .

```
public boolean isEmpty()
```

Variable

- The name should be Camel Case .

```
private long serialVersion = 1234L ;
```

Constants

- The name should be all uppercase letter. Compound words should be separated by underscores.

```
private int DEFAULT_CAPACITY = 10 ;
```

Enum

- Enum set names should be all uppercase letters.

```
public enum Duration {  
    SECOND , MINUTE , HOUR  
}
```

Acronyms

- Even though acronyms are generally represented by all Upper Case but in Java, only the first letter of acronyms should be upper case and rest lower case.

```
public void parseXml(){}
```

Questions

- What is the meaning of naming convention?
- Why do we need a naming convention?
- What is the difference between Camel Case and Pascal Case?
- What is the difference between Upper Camel Case and Lower Camel Case?
- Explain naming convention for Interface, Class, Method, Variable, Constant, Enum, and Acronyms?

OBJECT CLASS

Every Java class is inherited, directly or indirectly, from `java.lang.Object` class, which also means that a variable of `Object` class can reference the object of any class.

Because of inheritance, following `Object` class methods are available for overriding with class-specific code.

- `hashCode()` - returns hash-code value for the object.
- `equals()` - compares two objects for equality using identity (`==`) operator.
- `clone()` - creates copy of object. An `Overriding` class should inherit the `Cloneable` interface and implement a `clone()` method to define the meaning of copy.
- `toString()` - returns string representation of the object.
- `finalize()` - called by Garbage Collector to clean up resources. `Object`'s implementation of `finalize()` does nothing.

Questions

- What is the base class for all Java classes?
- What are the different methods of `Object` class available for overriding in the derived class?
- What happens if your class does not override `equals` method from the `Object` class?
 - *The `equals()` method in `Object` class compares whether object references are the same and not the content. To*

compare content, you need to override the equals() method.

- What is the purpose of the [clone](#) () method?
- Why should the overriding class define its implementation of the clone() method?
- What happens if the overriding class does not override the clone() method?
 - *In case if an object contains references to an external object, any change made to the referenced object will be visible in the cloned too.*

ACCESS MODIFIERS

The access modifiers determine the rules about whether other classes can access a variable or invoke a method.

At the class level, you can either use a [public](#) modifier or no modifier.

For class members, you can use one of the following access modifiers.

[private](#) - External classes cannot access the member.

[protected](#) - Only sub-classes can access the member.

[public](#) - All classes in the application can access the member.

no modifier - All classes within the package can access this member.

The access modifier in the overriding methods should be the same or less restrictive than the overridden method.

Optional [static](#) and [final](#) keywords are frequently used along with the access modifiers.

Questions

- What is the purpose of the access modifier?
- Is there any difference between the list of access modifiers available for a class and for its members?
- What is the scope of private, protected, and public access modifiers?
- What happens when no access modifier is specified with the class?

- If sub-class exists in a different package, can it still have visibility to the protected members of the super-class?
- Why should the member access modifier in the derived class be less restrictive than the base?
 - *As per the concepts of Inheritance, you should be able to use sub-class objects with super-class reference. This will not be possible if sub-class member is decelerated with more restrictive access modifier.*
- What should be the criteria to decide an access modifier for your class?
 - *You should use the most restrictive access modifier to ensure security and to prevent any misuse.*

STATIC

static class

Only nested/inner classes can be defined as `static` and not the outer class.

static variable and method

When a `static` keyword is used with the variables and the methods, it signifies that these members belong to the class and these members are shared by all the objects of the class. Static members do not have a copy and are stored only at a single location in memory. These members should be accessed using the class name.

A static method does not have access to instance methods or properties, because static members belong to the class and not the class instances.

Questions

- What are the static classes?
- Can any class be declared as a static class?
- What are the static methods?
- What are the static variables?
- Who owns the static class members? How is that different for non-static classes?
- How should you access class members that are static?
- Can a static method have access to an instance member? Why?

FINAL

final Class

A **final** class cannot be extended, which makes the class secure and efficient.

final Method

A **final** method cannot be overridden, which prevents any possibility of introducing any unexpected behavior to the class.

final Variable

A **final** variable reference cannot be changed, but the content of the [mutable object](#), that the final variable is referencing, can be changed.

blank final variable – a variable that is not initialized at the point of declaration.

Notes

- A **blank final** variable needs to be initialized in the constructor of the class.
- **The final** variables are like immutable variables, so computations related to final variables can be cached for optimization.

Questions

- Explain the final class?
- What are the benefits of declaring a class final?
- Explain the final method?

- What are the benefits of declaring a method final?
- Explain the final variable?
- What are the benefits of declaring a variable final?
- When you declare a variable final, can you change the content of the object it's referencing?
- When you declare a variable final, can you change it to reference another object?
- What is the blank final variable?
- How does declaring a variable as final helps with optimization?

STATIC INITIALIZATION BLOCK

- A [static initialization](#) block is generally used to ensure that all the required class resources (like drivers, connection strings, etc.) are available before any object of the class is used.
- A [static block](#) does not have access to the instance members.
- A [static block](#) is called only once for a class.
- A class can define any number of initializing blocks, which gets called in order of their definition in the class.
- You can only throw an unchecked exception from a [static block](#).

In this code example, [static initialization block](#) creates connection string once for the class.

```
private static String connectionString ;  
static {  
    connectionString = getConnectionSting();  
}
```

Questions

- What is a static initialization block?
- What is the primary purpose of a static initialization block? What are the types of things you should do in a static block?
- Can you access instance members from static initialization block? Why?

- Does the static initialization block gets called whenever an instance of the class is created?
- How many static blocks can be defined in a class?
- When multiple static blocks are defined, what is the criterion about their order of execution?
- Can you throw an exception from the static initialization block?
What type?

FINALLY

The primary purpose of a `finally` block is to ensure that the application is brought back to a consistent state, after the operations performed in the `try` block. Within the `finally` block, usually, the resources like streams and database connections can be closed to prevent leaks.

```
InputStream is = null ;
try {
    is = new FileInputStream("input.txt" );
}
finally {
    if (is != null ) {
        is.close();
    }
}
```

`finally` block execution

The compiler does all in its power to execute the `finally` block, except in the following conditions:

- If `System.exit()` is called.
- If the current thread is interrupted.
- If JVM crashes.

Return from `finally`

You must never return from within the `finally` block. If there is a `return` statement present in the `finally` block, it will immediately return, ignoring any other return present in the function.

Questions

- How do you guarantee that a block of code is always executed?

- What are the things that you should do in a finally block?
- What are the things that you should do in a catch block?
- Does finally block always execute? What are the conditions when finally block does not execute?
- Should you ever return from the finally block? Why?

FINALIZE()

When the [Garbage Collector](#) determines that there is no reference to an object exists, it calls `finalize()` on that object; just before removing that object from memory.

The `finalize()` method will not be called if an object does not become eligible for garbage collection, or if JVM stops before garbage collector gets a chance to run.

The `finalize()` method could be overridden to release the resources like file handles, database connections, etc.; but you must not rely on *finalize* to do so, and release such resources explicitly.

There is no guarantee that `finalize()` will be called by the [JVM](#), and you should treat `finalize` only as a backup mechanism for releasing resources. Where possible, use the [try-with-resource](#) construct to automatically release the resources; as soon as their usage is finished.

If an uncaught exception is thrown by the `finalize()` method, the exception is ignored before terminating the finalization.

Questions

- What is the `finalize` method in Java?
- When does the `finalize` method gets called?
- Who calls the `finalize` method?
- What are the things that can be done in the `finalize` method?
- Should you explicitly call the `finalize` method to release resources? Why?

- What are some alternate mechanisms that can be used to release system resources?
- What happens if an unhandled exception is thrown from the finalize method?

WIDENING VS NARROWING CONVERSIONS

Widening Conversions

Widening conversions deals with assigning an object of sub-class (derived class) to an object of super-class (base class). In the example below, `Car` is derived from `Vehicle` .

```
Car car = new Car ();  
Vehicle vehicle = car ;
```

Narrowing Conversions

Narrowing conversions deals with assigning an object of super-class (base class) to an object of sub-class (derived class). An explicit cast is required for conversion. In the example below, `Bike` is derived from `Vehicle` .

```
Vehicle vehicle = new Vehicle ();  
Bike bike = ( Bike ) vehicle ;
```

Questions

- What is widening conversion?
- What is narrowing conversion?
- Is there any possibility of a loss of data in narrowing conversion?

GETTERS AND SETTERS

The following code demonstrates the usage of [getter](#) and [setter](#) .

```
public class Person {  
  
    private String name ;  
  
    public String getName() {  
        return StringUtils.capitalize( name );  
    }  
  
    public void setName( String name) {  
        if (name.isEmpty()){  
            System . out .println( "Name string is empty" );  
            //throw exception  
        }  
        this . name = name;  
    }  
}
```

Benefits of using getter and setter

- Validations can be performed in the setter or can be added later when required.
- Value can have alternative representation, based on internal (storage) or external (caller's) requirement.
- It hides the internal data structure used to store the value.
- Internal fields can be changed, without requiring changing any users of the code.
- It encapsulates the internal complexity while retrieving or calculating the value.

- It provides the ability to specify different access modifiers for getter and setter.
- It provides the ability to add debugging information.
- It can be passed around as Lambda expressions.
- Many libraries like mocking, serialization, etc. expect getters/setters for operating on the objects.

Questions

- Why do you need getters and setters when you can directly expose the class field?
- Explain a few benefits of using getters and setters?

VARARGS VS OBJECT ARRAY

The **varargs** argument allows zero or more parameters to be passed to the method; whereas, an **object array** argument cannot be called with zero parameters.

varargs

```
public static int getCumulativeValue( int ... values){  
    int sum = 0 ;  
    for ( int value : values){  
        sum += value ;  
    }  
    return sum ;  
}
```

object array

```
public static int getCumulativeValues( int [] values){  
    int sum = 0 ;  
    for ( int value : values){  
        sum += value ;  
    }  
    return sum ;  
}
```

- The **varargs** can only be the last argument in the method; whereas, an object array can be defined in any order.
- Both **varargs** and object array are handled as an array within a method.
- Though **varargs** are not very popular, it can be used in any place where you have to deal with an indeterminate number of

parameters.

Questions

- What is varargs?
- What is the difference between varargs and object array?
- Can you call a method with zero arguments, that expects a varargs as a parameter?
- Can you overload a method that takes an int array, to take an int varargs?
- What are the different scenarios where you can use varargs?

DEFAULT INTERFACE METHOD

- **Default interface methods** are directly added to an Interface to extend its capabilities.
- The **Default interface method** can be added to an Interface that is not even under your control.
- **Default interface method** does not break any existing implementation of an interface it is added to.
- Implementing class can override the default methods defined in the interface.
- The **Default method** is also known as the **Defender** or **Virtual extension method**.

In this code example **default Interface method** , **getAdditonSymbol()** , is added to an existing interface **Calculator** .

```
public interface Calculator {  
    public < T > T  add( T num1, T num2);  
    default public String getAdditionSymbol(){  
        return "+" ;  
    }  
}
```

Limitations with the Default method

- If the class inherits multiple interfaces having default methods with the same signature, then the implementing class has to provide the implementation for that default method.
- If any class in the inheritance hierarchy has a method with the same signature, then default methods become irrelevant.

Default method vs Abstract method

Following are a couple of minor differences:

- Abstract methods allow defining constructor.
- Abstract methods can have a state associated.

With Default method - Abstract class vs Interface

With the introduction of **default methods** , now even the Interfaces can be extended to add more capabilities, without breaking the classes that inherit from the Interface.

Questions

- What are default interface methods?
- What are the benefits of default interface methods?
- Can you add default interface methods to an interface that is not directly under your control?
- Can you override the default interface methods to provide different implementation?
- What happens when a class inherits two interfaces and both define a default method with the same signature?
- How defining a default method in an interface is different from defining the same method in an abstract class?

STATIC INTERFACE METHOD

- **Static Interface methods** are directly added to an interface to extend its capabilities.
- **Static Interface methods** are generally used to implement utility functions like validations, sorting, etc.
- **Static interface methods** are also used when you want to enforce specific behavior in the classes inheriting the Interface.
- Implementing class cannot override the static methods defined in the interface it is inheriting.
- The **static Interface method** can even be added to an interface that is not under your control.
- Similar to the default Interface method, even the static interface method does not break any existing implementation of the interface it is added to.

In this code example static Interface method, `getUtcZonedDateTime()`, is added to an existing interface `DBWrapper`.

```
public interface DBWrapper {  
    static ZonedDateTime getUTCZonedDateTime(  
        Instant date ){  
        ZoneId zoneId =  
            TimeZone.getTimeZone( "UTC" ).toZoneId();  
        ZonedDateTime zonedDateTime =  
            ZonedDateTime.ofInstant(date, zoneId );  
        return zonedDateTime ;  
    }  
}
```

Questions

- What is static interface method?
- Where can you use static interface methods?
- Can you override static interface methods?
- What is the difference between static and default interface methods?
- Can you add a static interface method to an interface which is not directly under your control?
- What happens if a class inherits two interfaces and both define a static interface method with the same signature?

ANNOTATIONS

An [annotation](#) associates metadata to different program elements. Annotations may be directed at the compiler or at runtime processing.

[Annotations](#) metadata can be used for documentation, generating boilerplate code, performing compiler validation, runtime processing, etc. [Annotations](#) do not have any direct effect on the code piece they annotate.

We can apply annotations to a field, a variable, a method, a parameter, a class, an interface, an enum, a package, an annotation itself, etc.

Usage

User-defined annotations are directly placed before the item to be annotated.

```
@Length (max= 10 , min= 5 )  
public class ParkingSlot {  
    // Code goes here  
}
```

Built-in annotations

- [@Deprecated](#) - signifies that the method is obsoleted.
- [@Override](#) - signifies that a superclass method is overridden.
- [@SuppressWarnings](#) - used to suppress warnings.

Questions

- What are annotations?
- Where can you use annotations?

- Where you can apply annotations?

PREFERENCES

In Java, the [Preferences](#) class is used for storing user preferences in hierarchical form. [Preferences](#) class abstracts out the process of storage and stores the preferences in an operating system specific way: preferences file on Mac, or the registry on Windows systems. Though the keys in preferences are [Strings](#) value can belong to any primitive type.

Applications use Preferences class to store and retrieve user and system preferences and configuration data.

Questions

- What are the Preferences?
- What are the types of information that can be stored with the Preferences?
- While using the Preference class, should you worry about the internal format used by the operating system?

PASS BY VALUE OR PASS BY REFERENCE

In Java - method arguments, whether primitive or object reference, are always passed by value, and access to the objects is allowed only through the reference. While passing an object to a method, it's the copy of the reference that is passed and not the object itself. Any changes done to the object reference changes the object content and not the value of the reference.

Questions

- What is the difference between pass-by-value and pass-by-reference?
- How is the reference type argument passed in Java; by reference or by value?
- If a copy of the reference is passed by value, how can the method get access to the object that reference is pointing to?
- If a copy of the reference is passed by value, can you change the value of reference?

OBJECT ORIENTED PROGRAMMING

POLYMORPHISM

Polymorphism is an ability of a class instance to take different forms based on the instance it's acting upon.

Polymorphism is primarily achieved by sub-classing or by implementing an interface. The derived classes can have their own unique implementation for a certain feature and yet share some of the functionality through inheritance.

The behavior of objects depends specifically on its position in the class hierarchy.

Consider you have a **Furniture** class that has **addLegs()** method. A **Chair** and a **Table** class extend **Furniture** class; having their implementation of **addLegs()**. In the above situation, the runtime determines which implementation of **addLegs()** method gets called, based on whether you have a **Chair** or a **Table** class instance.

```
public abstract class Furniture {  
    public abstract void addLegs();  
    public void print( String message){  
        System . out .println(message);  
    }  
}
```

```
class Chair extends Furniture {  
    @Override  
    public void addLegs() {  
        print( "Chair Legs Added" );  
    }  
}
```



```
class Table extends Furniture{  
    @Override  
    public void addLegs() {  
        print( "Table Legs Added" );  
    }  
}
```

```
Furniture furniture = new Chair ();  
// This prints "Chair Legs Added"  
furniture .addLegs();
```

```
furniture = new Table ();  
// This prints "Table Legs Added"  
furniture .addLegs();
```

Benefits of polymorphism

The real power and benefit of polymorphism can be achieved when you can code to an abstract base class or an interface. Based on the context, **polymorphism** enables the selection of most appropriate class implementation. Not only in production code, but it also paves the way to have an alternate implementation for testing.

Questions

- What is Polymorphism?
- What are the different ways to achieve polymorphism?
- How is inheritance useful to achieve polymorphism?
- What are the benefits of polymorphism?
- How is the polymorphism concept useful for unit testing?

Parametric polymorphism

In Java, Generics provides an implementation of **Parametric Polymorphism**, which enables the use of the same code implementation with the values of different types, without compromising on compile-time type safety check.

In the example below, we added an upper bound to type parameter T such that it implements an interface that guarantees `getWheelsCount()` method in the type T.

```
interface Vehicle {  
    int getWheelsCount();  
}  
  
class Car < T extends Vehicle> {  
    private T vehicle ;  
    public Car( T vehicle) {  
        this . vehicle = vehicle;  
    }  
    public int getWheelsCount() {  
        return vehicle .getWheelsCount();  
    }  
}
```

It can take a list of **Vehicle** of type T and returns the count of wheels in the **Vehicle**, without worrying about what type T actually is.

Questions

- What is Parametric Polymorphism?
- How *Generics* is used to achieve Parametric Polymorphism?
- How are the *Type Wildcards* used to achieve Parametric Polymorphism?
- Can you achieve Parametric Polymorphism without Generics?

Subtype polymorphism

In **Subtype polymorphism**, also known as **inclusion polymorphism**, the parameter definition of a function supports any argument of a type or its subtype.

So if the parameters of a function have subtypes, then that function exhibits **subtype polymorphism**. Java code below illustrates the use of this kind of polymorphism.

```
abstract class Vehicle{  
    public abstract int getWheelsCount();  
}
```

```
class Car extends Vehicle{  
    @Override  
    public int getWheelsCount() {  
        return 4 ;  
    }  
}
```

```
class Bike extends Vehicle{  
    @Override  
    public int getWheelsCount() {  
        return 2 ;  
    }  
}
```

```
public void printWheelsCount(Vehicle vehicle) {  
    print(vehicle.getWheelsCount());  
}
```

```
public void main( String [] args) {  
    printWheelsCount( new Car ());  
}
```

```
    printWheelsCount( new Bike ());  
}
```

In the code above, the method `printWheelsCount()` takes a `Vehicle` as parameter and prints count of wheels in the `Vehicle`. The main method shows subtype polymorphic calls, taking objects of `Car` and `Bike` as arguments. Any place that expects a type as a parameter will also accept subclass of that type as a parameter.

Questions

- What is Subtype Polymorphism?
- What is Inclusion Polymorphism?
- What is the difference between Parametric Polymorphism and SubType Polymorphism?
- Can you achieve SubType polymorphism using Generics?

OVERRIDING

- Method **overriding** is redefining the base class method to behave in a different manner than its implementation in the base class.
- Method **overriding** is an example of dynamic or runtime polymorphism.
- In **dynamic polymorphism** , the runtime takes the decision to call an implementation, as the compiler does not know what to bind at compile time.

Rules for method overriding

- Method arguments and their order must be the same in the overriding method.
- Overriding methods can have the same return type or subtype of the base class method's return type.
- The access modifier of the overridden method cannot be more restrictive than its definition in the base class.
- Constructor, **static**, and **final** method cannot be overridden.
- The overridden method cannot throw **checked exceptions** if its definition in base class doesn't, though the overridden method can still throw an **unchecked exception** .

Questions

- What is method overriding?
- What is dynamic polymorphism?
- Why can't you override static methods defined in super-class or interface?

- Can you override a final method defined in super-class?
- Can you override a public method in super-class and mark it protected?
- Why can't you override constructor of super-class?
- Can an overridden method throw checked exception; when the method in super-class does not? Why?
- What are the benefits of method overriding?

@Override

The `@Override` annotation is a way to explicitly declare the intention of overriding a method implementation defined in the base class. Java performs compile-time checking for all such annotated methods. It provides an easy way to mistake-proof against accidentally writing the wrong method signature when you want to override from the base class.

If a derived class defines a method having the same signature as a method in the base class, the method in the derived class hides the one in the base class. By prefixing a subclass's method header with the `@Override` annotation, you can detect if an inadvertent attempt is made to overload instead of overriding a method.

Questions

- What is the purpose of the `@Override` annotation?
- What happens if you define a method with the same signature as defined in the super-class and without using the `@Override` annotation?
- What are the benefits of `@Override` annotation?

OVERLOADING

- The method **overloading** is about defining more than one method with the same name, but with different parameters.
- Method **overloading** is an example of static or compile-time polymorphism.
- In **static polymorphism** , it's while writing the code the decision is made to call a specific implementation.

Rules for method overloading

- The method can be overloaded by defining method with the same name as an existing one but having
 - A different number of arguments.
 - The different datatype of arguments.
 - A different order of arguments.
- The return type of the overloaded method can be different.
- A method with the same name and the same parameter cannot be defined when they differ only by return type.
- Overloading methods are not required to throw the same exception as the methods it's overloading.

Operator Overloading

Operator overloading is an ability to enhance the definition of the language-dependent operators. For example, you can use **+** operator to add two integers and also to concat two strings.

Questions

- Explain method overloading?

- What is static polymorphism?
- What is the difference between static and dynamic polymorphism?
- Can you override a method such that all the parameters are the same with the difference only in return type?
- What is operator overloading?
- What are the benefits of method overloading?
- What is the difference between overriding and overloading?

ABSTRACTION

[Abstraction](#) helps to move the focus from the internal details of the concrete implementation to the types and its behavior. [Abstraction](#) is all about hiding details about data, its internal representation, and its implementation.

The other equally important object-oriented concept is an [encapsulation](#) that could be used to abstract the complexities and the internal implementation of a class.

[Abstraction](#) also helps to make the software maintainable, secure, and provides an ability to change the implementation without breaking any client.

[Questions](#)

- What is abstraction?
- How abstraction is different from encapsulation?
- What are the benefits of abstraction?
- Can you achieve abstraction without encapsulation?

INHERITANCE

[Inheritance](#) is an object-oriented design concept that deals with reusing an existing class definition (known as super-class) and defining more special categories of class (known as sub-class) by inheriting that class. It focuses on establishing the [IS-A](#) relationship between sub-class and its super-class. Inheritance is also used as a technique to implement [polymorphism](#) ; when a derived type implements a method defined in the base type.

Rules for Inheritance

- There can be a multiple levels of inheritance, based on the requirements to create specific categories.
- Only single class inheritance is allowed in Java, as multiple inheritances comes with its share of complexity; check [Diamond Problem](#) .
- A class declared [final](#) cannot be extended.
- A class method declared [final](#) cannot be overridden.
- Constructors and private members of the base class are not inherited.
- The constructor of the base class can be called using [super\(\)](#) .
- The base class's overridden method can be called using [super](#) keyword, otherwise, you will end up calling an overriding method recursively.

Questions

- Explain inheritance?
- What is the purpose of inheritance?

- What should be the criteria to decide the inheritance relation between two classes?
- How inheritance plays an important role in polymorphism?
- Can you inherit the final class?
- What happens if you don't use the super keyword to call an overridden member?
- Why can't you inherit static members defined in the super-class?
- What are the challenges one can face if multiple inheritances are possible in Java?

COMPOSITION

The **composition** is an object-oriented design concept that is closely related to **inheritance** , as it also deals with reusing classes; but it focuses on establishing a **HAS- A** relationship between classes. So unlike **Inheritance** , which deals with extending features of a class, composition reuses a class by composing it. The **composition** is achieved by storing reference of another class as a member.

Inheritance vs Composition

Primary issue with **inheritance** is that it breaks **encapsulation** as the derived class becomes tightly coupled to implementation of the base class. The problem becomes complex when a class is not designed keeping future inheritance scope and you have no control over the base class. There is possibility of breaking a derived class because of changes in the base class.

So, **inheritance** must be used only when there is perfect **IS-A** relationship between the base and the derived class definitions; and in case of any confusion prefer **composition** over **inheritance** .

Questions

- Explain composition?
- What is the difference between inheritance and composition?
- What should be the criteria to decide composition relation between two classes?
- Explain few problems with inheritance that can be avoided by using composition?

- When should you prefer composition over inheritance and vice versa?

FUNDAMENTAL DESIGN CONCEPTS

DEPENDENCY INJECTION VS INVERSION OF CONTROL

The [Dependency Injection](#) and the [Inversion of Control](#) promotes modular software development by loosely coupling the dependencies. Independent modular objects are also more maintainable and testable.


[Inversion of Control](#)

[Inversion of Control](#) provides a design paradigm where dependencies are not explicitly created by the objects that require these, but such objects are created and provided by the external source.

[Dependency Injection](#)

[Dependency Injection](#) is a form of IoC that deals with providing object dependencies runtime; through constructors, setters or the service locators. Annotations and Interfaces are used to identify the dependency sources to create and inject dependencies.

- Mode of dependency injection:
 - Through constructor
 - Through setter
 - Through method parameter
- It's the responsibility of the dependency injection framework to inject the dependencies.



The code below demonstrates dependency injection as a constructor parameter.

```
public class Account {  
    UserService userService ;  
    AccountService accountService ;  
  
    public Account(UserService userService,  
                   AccountService accountService) {  
        this.userService =  
            userService;  
        this.accountService =  
            accountService;  
    }  
}
```

Not only in production systems, but **DI** and **IoC** also provide immense help in unit testing too, by providing an ability to mock dependencies. **Spring framework** is an example of a DI container.

Note

It is important to ensure that dependency objects are initialized before they are requested.


Questions

- What is Inversion of Control?
- What is Dependency Injection?
- What is the difference between the Inversion of Control and Dependency Injection?
- What are the different ways to implement Dependency Injection?
- What the different ways to identify dependency sources?
- Who has the responsibility to inject dependent objects?
- How Dependency injection is useful for unit testing?
- How Dependency Injection can be used for modular software development?

SERVICE LOCATOR

The [service locator](#) is an object that encapsulates the logic to resolve the service requested. The [service locator](#) also provides an interface to register services with it, which allows you to replace the concrete implementation without modifying the objects that depend on these services.

In the figure below, [Account](#) class uses [ServiceLocator](#) to resolve the [Account Service](#) and [User Service](#) it depends on.



```
public class Account {  
    UserService userService ;  
    AccountService accountService ;  
  
    public Account() {  
        this.userService =  
            ServiceLocator.getService(UserService.class );  
        this.accountService =  
            ServiceLocator.getService(AccountService.class );  
    }  
}
```

Benefits of Service Locator

- The class does not have to manage any service dependency and its life cycle.
- It helps to test a class independently, without the availability of real services class depends on.
- It enables runtime optimization; as services can be registered and unregistered runtime.


Questions

- Explain the Service Locator design pattern?
- What are the benefits of using a Service Locator design pattern?
- What is the difference between Service Locator and Dependency Injection pattern?
- When would you prefer Service Locator over Dependency Injection and vice versa?
- How does Service Locator help in testing?

DIAMOND PROBLEM

Java doesn't allow extending multiple classes because of the ambiguity that could arise when more than one super-classes have the same signature method, and the compiler can't decide which super-class method to use.

Consider the inheritance hierarchy set out in the figure below. If the method `calculate()` defined in the `Base` class is overridden by both, `DerivedLeft` and `DerivedRight`, then it creates ambiguity regarding which version of `calculate()` does the `Confused` class inherits.



In the code below there is an ambiguity regarding which version of `calculate()` should be called. This is known as **Diamond Problem** in Java.

```
public static void main (String [] args){  
    Base base = new Confused ();
```

```
    base.calculate();  
}
```

Diamond Problem with Default Interface Method

Diamond problem is possible with the introduction of Default Interface methods. If `Base`, `DerivedLeft`, and `DerivedRight` are Interfaces, and there exists `calculate()` as default interface method in all three, it will cause the Diamond Problem.

In such a scenario the `Confused` class has to explicitly re-implement the `calculate()` method; otherwise, the ambiguity will be rejected by the compiler.

Questions

- Explain the Diamond Problem in Java?
- Why Java does not provide multiple inheritances?
- Using default interface methods, a class can still inherit two interfaces with the same signature method; would this not cause Diamond Problem? How can you solve it?

PROGRAMMING TO INTERFACE

[Programming to interface](#) forms the basis for modular software development by facilitating decoupling between software components. A high level of decoupling improves maintainability, extensibility, and testability of software components. Modular software design also helps to improve speed to market, as it facilitates parallel software development between multiple teams working with the same code base.

It's the [Programming to Interface](#) design paradigm that forms the foundation for [Inversion of Control](#), which manages dependency relationships in any large software application.

Let take a very simple example. Suppose we have a method to sort a collection which accepts [Map](#) Interface as the parameter. It means that the [sort\(\)](#) method is not tied to any specific type of [Map](#) implementation and you can pass any concrete implementation of [Map](#) interface.

```
public static void main ( String [] args){  
    sort( new HashMap<>());  
    sort( new TreeMap<>());  
    sort( new ConcurrentSkipListMap<>());  
    sort( new TreeMap<>());  
}
```

```
public static void sort(Map map){  
    // perform sort  
}
```

[Benefits of programming to interface](#)

- Based on the context, you can runtime select the most appropriate behavior.
- For testing, you can pass mock objects or stubs implementation.
- The interface/API definitions or the contract does not change frequently.
- Programming to Interface also facilitates parallel development between teams, as developers from different team can continue writing code against an interface before doing final integration.

Questions

- What is the concept of programming to interface?
- What are the benefits of programming to interface?
- How does programming to interface facilitate decoupling between software components?
- How dependency injection and programming to the interface are inter-related? Can you achieve dependency injection without supporting programming to interface?
- What are the benefits of modular software?
- How does programming to interface help with unit testing?

ABSTRACT CLASS VS INTERFACE

Abstract Class

An [Abstract class](#) cannot be instantiated but can be extended. Extend abstract class when you want to enforce a common design and implementation among classes, and expect the derived class to implement required specialized functionality.

Interface

The [Interface](#) is a set of related methods, which defines its behavior and its contract with the outside world. Use an interface when you want to define common behavior among unrelated classes. An interface can also be defined without any method and such interface is known as [marker interface](#); such interfaces are generally used to categorize the classes. An example of a marker interface is [java.io.Serializable](#) , which does not define any method but must be implemented by the classes that support serialization.

Difference between Abstract Class and Interface

- An Abstract class can be updated to add more capabilities to a class whereas Interface can be added to implement new behavior to the class. Though with the introduction of [default interface methods](#) , even Interfaces can be extended to have more capabilities.
- An Interface can be multiple inherited; whereas, an abstract class cannot.
- Interface define behavior, so it can even be applied to unrelated classes; whereas, related classes extend Abstract class.

- Abstract class methods can have any type of access modifier; whereas, Interface has all public members.
- An Abstract class can have state associated, which is not possible with Interface.
- An Abstract class can be extended without breaking the classes that extend it; whereas, any change in Interface, except for non-default and non-static methods, will break the existing implementation.

Questions

- If an abstract class cannot be instantiated, then what could be a need to define a constructor for an Abstract class?
 - *A constructor can be used to perform the required field initialization and also to enforce class constraints.*
- Define Abstract class? What role an Abstract class plays in a class design?
- Define Interface? What role does Interface play in class design?
- When would you prefer using Abstract class over Interface and vice-versa?
- Explain various differences between Abstract Class and Interface?
- What are marker interfaces? How are marker interfaces used?
- Can you declare an interface method static?
- With the introduction of default interface methods; how Abstract class is still different from an Interface?

INTERNATIONALIZATION AND LOCALIZATION

Internationalization

The internationalization of software is the process to ensure that software is not tied to only one language or locale. Its shortened name is **i18n** .

Localization

The localization of software is the process to ensure that software has all the resources available to support a specific language or locale. Its shortened name is **l10n** .

Note

Internationalization facilitates localization.

Questions

- What is Internationalization?
- What is localization?
- What is the difference between localization and internationalization?
- Can you achieve localization without building support for Internationalization?

IMMUTABLE OBJECTS

An object is considered **immutable** when there is no possibility of its state change after its construction.

Advantages

- It is easier to design and implement an immutable object, as you don't have to manage state change.
- Immutable objects are inherently thread-safe because they cannot be modified after creation. So there is no It is easier to design and implement an immutable object, as you don't have to manage state change.
- Immutable objects are inherently thread-safe because they cannot be modified after creation. So there is no need to synchronize access to it.
- Immutable objects reduce Garbage Collection overhead.
- need to synchronize access to it.
- Immutable objects reduce Garbage Collection overhead.

Disadvantages

- A separate object needs to be defined for each distinct value because you cannot reuse an Immutable object.

The rule for defining Immutable Objects

- Declare the class **final** .
- Allow only the constructor to create the object. Don't provide field setter.
- Mark all the fields **private** .

Example of an immutable class, *Employee* .

```
final public class Employee {  
  
    final private int id ;  
    final private String name ;  
    final private String department ;  
  
    public Employee( int id,  
                    String name,  
                    String department) {  
        this . id = id;  
        this . name = name;  
        this . department = department;  
    }  
  
    public int getId() {  
        return id ;  
    }  
  
    public String getName() {  
        return name ;  
    }  
  
    public String getDepartment() {  
        return department ;  
    }  
}
```

Questions

- What is an immutable object?
- What are the rules for defining an immutable object?
- What are the advantages/disadvantages of an immutable object?
- How do you create an immutable object?

- What are the different situations where you can use immutable objects?
- What is the difference between a final and an immutable object?
- What is the optimization benefit of declaring a variable final?
- Can you list some of the problems with Immutability?
 - *It's harder to define constructors with loads of arguments.*
 - *Since the enforcement of immutability is left to the developer, even a single setter added accidentally can break it.*

CLONING

Cloning is the process of creating a copy of an object.

Simply assigning an existing object reference to an object, results in two references pointing to the same object.

There are two types of cloning, **shallow** cloning and **deep** cloning.

Shallow Cloning

Shallow cloning simply copies the values of the properties; for primitive property members, an exact copy is created; and for reference type members, its address is copied. So, for the reference type members, both original and the newly created copy, point to the same object in heap.

Deep Cloning

Deep cloning recursively copies the content of each member to the new object. Deep cloning always creates an independent copy of the original object. To create a deep clone, a dedicated method known, as **CopyConstructor** should be written.

Questions

- What is cloning?
- What is shallow cloning?
- Explain drawbacks with shallow cloning?
- What is deep cloning?
- What is CopyConstructor?
- When would you prefer deep cloning over shallow cloning and vice versa?

DATA TYPES

NAN

Not a Number also is known as NaN, is the undefined result produced because of arithmetic computations like divide by zero, operating with infinity, etc. No two NaNs are equal.

NaNs are of two types:

Quiet NaN - When a quiet NaN is results, there is no indication unless the result is checked.

Signalling NaN - When a signaling NaN results, it signals invalid operation expression.

Questions

- What is NaN or Not a Number?
- What is Quiet NaN?
- What is Signalling NaN?
- Are two NaNs equal?

ENUMSET

- The `EnumSet` is a specialized set implementation to be used with an Enum type.
- The `EnumSet` is represented internally as bit vectors, in a very compact and efficient manner.
- The `EnumSet` provides optimized implementation to perform bit flag operations and should be used in place of performing int flag operations.

The following code demonstrates the usage of EnumSet.

```
private enum Vehicle {  
    CAR ,  
    JEEP ,  
    MOTORCYCLE ,  
    SCOOTER  
};  
  
public static void main( String [] args){  
    EnumSet< Vehicle > TWOWHEELERS =  
        EnumSet.of( Vehicle . MOTORCYCLE ,  
                    Vehicle . SCOOTER );  
    if ( TWOWHEELERS .contains( Vehicle . MOTORCYCLE ){  
    }  
}
```

Questions

- What is EnumSet?

- Why should you prefer EnumSet to perform int based bit flag operations?

COMPARING THE TYPES

Primitive types can be compared only by using the [equality operators](#) (== and !=); whereas, reference types can be compared using both [equality operator](#) and [equals\(\)](#) method, depending upon what we want to compare.

[Equality Operators](#)

For reference types, the equality operators == and != are used to compare the addresses of two objects in memory and not their actual content.

[.equals\(\) Method](#)

Use [equals\(\)](#) method when you want to compare the content of two objects. [Object](#) class in Java defines [equals\(\)](#) method, which must be overridden by the subclasses to facilitate content comparison between its objects. If the [equal\(\)](#) method is not overridden by a class, then the [equals\(\)](#) method of the [java.lang.Object](#) class is called, which uses an [equality operator](#) to compare references.

[Questions](#)

- What are the different ways to compare types in Java?
- What does the equality operator compare for the reference types?
- What does the equals method compare?
- When would you prefer using equals method to equality operator?
- What happens if a class does not override the equals method?

Float Comparison

Two float numbers should not be compared using equality operator `==`; as the same number, say 0.33, may not be stored in two floating-point variables exactly as 0.33, but as 0.3300000007 and 0.329999999767.

So to compare two float values, compare the absolute difference of two values against a range.

```
if ( Math.abs( floatValue1 - floatValue2 ) < EPSILON ){  
    //  
}
```

where `EPSILON` is a very small number like 0.0000001, and that depend on the desired precision.

Questions

- Why shouldn't you use the equality operator to compare float value?
- What is the preferred way to compare two float values?
- What should be the criteria to decide on the value of the range that should be used to compare two float values?

String Comparison

`String` class object uses [equality operator](#) , `==`, to tests for reference equality and `equals()` method to test content equality.

You should always use `equals()` method to compare the equality of two `String` variables from different sources. [Interned](#) Strings can be compared using equality operators too.

Questions

- Why should you use the `equals` method to compare `String` objects?
- What is the pitfall of using an equality operator to compare two `String` objects?
- What are interned string? How can you compare two interned strings?

Enum Comparison

Enum can neither be instantiated nor copied. So always only a single instance of the enum is available, the one defined with the enum definition.

As only one instance of the enum is available, you can use both equality operator (==) and equals() method for comparison. But prefer using equality operator, ==, as it does not throw **NullPointerException** and it also performs compile-time compatibility check.

Questions

- What are the different ways to compare two enums?
- Explain why you can use both equality operator and equals method to compare enum?
- Which is the preferred way to compare two enum values?

ENUM VS PUBLIC STATIC FIELD

The code below demonstrates the usage of `enum` and `public static` field.

The following are advantages of using enum over public static int.

- The Enums are compile-time checked, whereas int values are not. With `public static int`, you can pass any `int` value to `AddVehicle()` method.
- An `int` value needs to be validated against an expected range; whereas, enums are not.
- Bitwise flag operations are built into `enumSet`.

`enum`

```
private enum Vehicle {  
    CAR ,  
    JEEP ,  
    MOTORCYCLE ,  
    SCOOTER  
};  
  
public void AddVehicle( Vehicle vehicle){  
}
```

`public static field`

```
public static int CAR = 0 ;  
public static int JEEP = 1 ;  
public static int MOTORCYCLE = 2 ;  
public static int SCOOTER = 3 ;
```



```
public void AddVehicle( int vehicle){  
}
```

Questions

- What are the advantages of using enum over public static int fields?
- Why do you need to perform extra validation with the int parameter as compared to the enum parameter?
- Why should you prefer using an enum over a public static int?

WRAPPER CLASSES

Each Primitive data type has a class defined for it that wraps the primitive datatype into the object of that class. Wrapper classes provide lots of utility methods to operate on primitive data values. As the wrapper classes enable primitive types to convert into reference types, these can be used with the collections too.

Questions

- What are the wrapper classes?
- What are the advantages of using the wrapper type over the primitive type?
- How can you use primitive types with collections?

AUTO BOXING AND AUTO UNBOXING

[Auto boxing](#) is an automatic conversion of primitive type to an object, which involves dynamically memory allocation and initialization of corresponding Wrapper class object. [Auto unboxing](#) is the automatic conversion of a Wrapper class to a primitive type.

In the code below, value [23.456f](#) is auto boxed to an object [Float\(23.456f \)](#) and the value returned from [addTax\(\)](#) is auto unboxed to float.

```
public static void main( String [] args){  
    float beforeTax = 23.456f ;  
    float afterTax = addTax( beforeTax );  
}  
  
public static Float addTax( Float amount){  
    return amount * 1.2f ;  
}
```

Questions

- What are auto boxing and auto unboxing?
- What are the advantages of auto boxing?

BIGINTEGER AND BIGDECIMAL

`BigInteger` and `BigDecimal` are used to handle values that are larger than `Long.MAX_VALUE` and `Double.MAX_VALUE`. Such large values are passed as String values to the constructor of `BigInteger` and `BigDecimal`. `BigDecimal` supports utility methods to specify the required rounding and the scale to be applied.

```
BigInteger bInt =  
    new BigInteger( "9876543210987654321098765" );
```

Both `BigInteger` and `BigDecimal` objects are immutable, so any operation on it creates a new object. `BigInteger` is mainly useful in cryptographic and security applications.

Questions

- What are the `BigInteger` and `BigDecimal` types?
- How are the values of `BigInteger` and `BigDecimal` internally stored?
- What are the usages of `BigInteger`?

STRINGS


STRING IMMUTABILITY

The `String` object is immutable, which means once constructed, the object which `String` reference refers to, can never change. Though you can assign the same reference to another `String` object.

Consider the following example:

```
String greeting = "Happy" ;  
greeting = greeting + " Birthday" ;
```

The code above creates three different `String` objects, " Happy ", " Birthday " and " Happy Birthday ".



- Though you cannot change the value of the `String` object you can change the reference variable that is referring to the object. In the above example, the `String` reference `greeting` starts referring to the `String` object "Happy Birthday".
- Note that any operation performed on `String` results in the creation of a new `String`.
- `String` class is marked `final` , so it's not possible to override the immutable behavior of the `String` class.

Advantages

- As no synchronization needed for String objects, it's safe to share a [String](#) object between threads.
- As [String](#) does not change, the Java environment caches [String](#) literals into a special area in memory known as a [String Pool](#) for optimization. If a [String](#) literal already exists in the pool, the same string literal is shared.
- Immutable [String](#) values safeguard against any change in value during execution.
- As hash-code of [String](#) object does not change, it is possible to cache hash-code and not calculate every time it's required.

Disadvantages

- The [String](#) class cannot be extended to provide additional features.
- If loads of [String](#) literals are created, either as new objects or because of any string operation, it will put the load on Garbage Collector.

Questions

- Why String objects are called Immutable?
- How is the String object created in memory?
- What are the advantages/disadvantages of String Immutability?
- Why String objects are considered thread-safe?
- What are the advantages of declaring the String class final?
- What memory optimization is performed by the Java environment for strings?
- Why you don't have to re-calculate hash-code of String object over the time it's used?

STRING LITERAL VS OBJECT

String Literal

String literal is a Java language concept where the **String** class is optimized to cache all the Strings created within double quotes, into a special area known as **String Pool**.

```
String cityName = "London" ;
```

String Object

The **String object** is created using the **new()** operator into the heap, like any other object of reference type.

```
String cityName = new String ( "London" );
```

Questions

- What is String literal?
- What are the differences between a String Literal and a String Object?
- How are the String Literals stored?

STRING INTERNING

- **String interning** is a concept of storing only a single copy of each distinct immutable String value.
- When you define any new **String literal** it is **interned** . Same **String** constant in the pool is referred for any repeating String literal.
- **String pool** literals are defined not only at the compile-time but also during runtime. You can explicitly call a method **intern()** on the **String** object to add it to the **String Pool** , if not already present.
- Placing an extremely large amount of text in the memory pool can lead to memory leak or performance issues.

Note: Instead of using String object, prefer using a string literal so that the compiler can optimize it.

Questions

- What is String interning?
- Can String Objects also be interned?
- What happens when you store a new String literal value that is already present in the string pool?
- What are the drawbacks of creating a large number of String literals?
- Which one is preferred: String Object or String Literal? Why?


STRING POOL MEMORY MANAGEMENT

The [String pool](#) is a special area in memory managed by the Java compiler for String memory optimization. If there is already a String literal present in the string pool, compiler refers the new String literal reference to an existing String variable in the pool, instead of creating a new literal. Java compiler can perform this optimization because String is immutable.

In this example below, both the String objects are different and are stored into [Heap](#) .

```
String cityNameObj = new String( "London" );
```


```
String capitalObj = new String ( "London" );
```



Whereas in this example below, both String literal reference to the same object in the memory pool.

```
String cityName = "London" ;
```

```
String capital = "London" ;
```



Questions:

- Explain String Pool Memory Management?
- How are String Literals stored in memory?
- How String Pool is optimized for memory?
- How are String Objects stored in memory?
- Why can Java not use a String Pool-like mechanism to store objects of other types of data?

IMMUTABILITY - SECURITY ISSUE

It's the responsibility of the [Garbage Collector](#) to clear string objects from the memory; though you can also [use reflection](#) to do so, that's not recommended.

Since Strings are kept in [String Pool](#) for re-usability, chances are that the strings will remain in memory for a long duration. As String is [immutable](#) and its value cannot be changed, a memory dump or accidental logging of such String can reveal sensitive content like password or account number, stored into it.

So instead, it's advisable to use char array ([char \[\]](#)) to store such sensitive information, which can be explicitly overwritten by an overriding content, thus reducing the window of opportunity for an attack.

Questions:

- How are String literals cleared from the String Pool?
- Can you use reflection to clear a String object?
- What are the security issues associated with the immutable Strings?
- Why should you not use Strings to store sensitive information such as passwords, access keys, etc.?
- What other data type can be used to store the password instead of String?

CIRCUMVENT STRING IMMUTABILITY

Immutability feature in String can be bypassed with reflection, though it is NOT recommended to use reflection to do so because it is a violation of security and is considered an attack. The following code demonstrates how reflection can be used to circumvent string immutability:

```
String accountNo = "ABC123" ;

Field field = String . class .getDeclaredField( "value" );
field .setAccessible( true );

char [] value = ( char []) field .get( accountNo );

// Overwrite the content
value [ 0 ] = 'X' ;
value [ 1 ] = 'Y' ;
value [ 2 ] = 'Z' ;

// Print "XYZ123"
System . out .println( accountNo );
```

Questions

- Can you override the String class to modify its immutability?
- Is it possible to circumvent string immutability?
- Is it recommended to circumvent string immutability using reflection? Why?

STRINGBUILDER VS STRINGBUFFER

Similarities

- Both `StringBuilder` and `StringBuffer` objects are mutable, so allows String values to change.
- The objects of both the classes are created and stored in heap.
- Both the classes have similar methods.

Differences

- `StringBuffer` methods are `synchronized` , so its thread-safe whereas `StringBuilder` is not.
- The performance of `StringBuilder` is significantly better than `StringBuffer` , as `StringBuilder` does not have synchronization overheads.

Note: If you need to share String objects between threads then use `StringBuffer` , otherwise `StringBuilder` .

Questions

- What are the similarities and differences between `StringBuffer` and `StringBuilder`?
- When would you prefer `StringBuffer` to `StringBuilder`?
- Between `StringBuffer` and `StringBuilder`, which one would you prefer in a single-threaded application?

UNICODE

Unicode is an international standard character encoding system that represents most of the written languages in the world. Before Unicode, there were multiple encoding systems prevalent: ASCII, KOI8, ISO 8859, etc., each encoding system having its code values and character set with different lengths. So to solve this issue, a uniform standard is created, which is known as Unicode. Unicode provides a platform and language independent unique number for each character.

Questions

- What are Unicode characters?
- What are the advantages of using Unicode characters?
- What are the encoding systems problems which predominated before Unicode?

INNER CLASSES

INNER CLASSES

Inner Class – is a class within another class.

Outer Class – is an enclosing class that contains the inner class.

Note

- The compiler generates a separate class file for each inner class.

Advantages of inner classes

- It's easy to implement callbacks using inner classes.
- The inner class has access to the private members of the enclosing class that aren't even available to the inherited classes.
- The inner class helps to enforce closures that make the enclosing instance and the surrounding scope available.
- The outer classes provide additional namespace to the inner classes.

Questions

- What is the inner class?
- What is the outer class?
- What are the advantages of defining an inner class?
- What are closures? How inner class can be used to create closures?

- What are callbacks? How inner class can be used to create callbacks?
- Can an inner class access private members of the enclosing outer class?
- What benefit does the outer class bring?

STATIC MEMBER NESTED CLASS

- A **static nested class** is declared as static inside a class like any other member.
- A **static nested class** is independent and has nothing to do with the outer class. It is generally used to keep together with the outer class.
- It can be declared public, private, protected, or at the package level.

Declaration of static nested class

```
// outer class
public class Building {
    // static member inner class
    public static class Block{

    }
}
```

Creating an object of static nested class

```
// instance of static member inner class
Building.Block block =
    new Building.Block();
```

Questions

- What is a static nested class?

- Static classes, nested or outer, are the same; then what's the benefit of defining an inner class as static?

LOCAL INNER CLASS

- A **local inner class** is declared and can be used inside the method block.
- It cannot be declared public, private, protected, or at the package level.

Creation of local inner class

```
// outer class
public class CityNames {
    private List <String> cityNames =
        new ArrayList<>();

    public Iterator <String> nameIterator(){
        // local inner class
        class NameIterator
            implements Iterator <String> {
            @Override
            public boolean hasNext() {
                return false ;
            }
            @Override
            public String next() {
                return null ;
            }
        }
        // return an instance of a local inner class.
        return new NameIterator();
    }
}
```

Note

- To be able to use the inner class outside, the local inner class must implement a public interface or Inherit a public class and override methods to redefine some aspects.

Questions

- What is the difference between an inner class and a local inner class?
- Why can't you use an access modifier with the local inner class?
- Explain the rules for defining local inner class?
- What problem does a local inner class solves?

NON-STATIC NESTED CLASS

- A **non-static nested class** is declared inside a class like any other member.
- It can be declared public, private, protected, or at the package level.
- Non-static nested classes are closures as they have access to enclosing instance.
- An object of the outer class is required to create an object of a non-static inner class.

Declaration of non-static nested class

```
// outer class
public class Building {
    // non-static member inner class
    public class Block{
    }
}
```

Creating an object of non-static nested class

```
// instance of the outer class
Building building =
    new Building();
// instance of non-static member inner class
Building.Block block =
    building . new Block();
```

Questions

- What is a non-static nested class?
- Why a non-static nested class can be used as closures?
- Can you create an instance of a non-static inner class without defining an outer class?

ANONYMOUS INNER CLASS

- Anonymous inner class does not have a name.
- The anonymous inner class is defined and its object is created at the same time.
- Anonymous inner class is always created using `new` operator as part of an expression.
- To create an Anonymous class, a `new` operator is followed by an existing interface or class name.
- The anonymous class either implements the interface or inherits from an existing class.

Creation of anonymous inner class

```
// outer class
public class CityNames {
    private List <String> cityNames =
        new ArrayList<>();

    public Iterator <String> nameIterator(){
        // Anonymous inner class
        Iterator <String> nameIterator =
            new Iterator <String> () {
                @Override
                public boolean hasNext() {
                    return false ;
                }
            }
        @Override
        public String next() {
            return null ;
        }
    }
}
```

```
    }  
};  
// return an instance of a local inner class.  
return nameIterator();  
}  
}
```

Notes

- Do not return inner classes as it has reference to the outer enclosing class, otherwise memory leaks may occur .
- Use anonymous class when you want to prevent anyone from using the class anywhere else.
- Serialization of Anonymous and Inner classes must be avoided, as there could be compatibility issues during de-serialization, due to different JRE implementation.

Questions

- What is an anonymous inner class?
- How an anonymous inner class is different from the local inner class?
- Why you shouldn't return an inner class object from a method?
- If you want to prevent anyone from using your class outside, which type of inner class you should define?
- Why should you avoid serialization of anonymous and inner classes?

FUNCTIONAL PROGRAMMING

FUNCTIONAL INTERFACE

Functional Interface is an interface with only one abstract method but can have any number of default methods.

```
@FunctionalInterface
public interface Greater<T> {
    public T greater(T arg1, T args2);
}
```

Annotation **@FunctionalInterface** generates a compiler warning when the interface is not a valid functional interface.

```
@FunctionalInterface
public interface Greater<T> {
    public T greater(T arg1, T args2);
}
```

Account class, defined below, used as an argument to the functional interface **Greater** .

```
public class Account {
    private int balance ;

    public Account(int balance) {
        this .balance = balance;
    }

    public int getBalance() {
        return balance ;
    }

    @Override
    public String toString() {
        return "Account{" +
            "balance=" + balance +
            '}';
    }
}
```

The code below demonstrates the usage of [Lambda expression](#) to find the account with the greater balance. Similarly, the same [Function Interface](#) , [Greater](#) , can be used to compare other similar business objects too.

```
public static void main(String [] args){  
    Greator<Account > accountComparer =  
        (Account acc1, Account acc2) ->  
            acc1.getBalance() > acc2.getBalance() ?  
                acc1 :  
                acc2;  
  
    Account account1 =  
        new Account(6 );  
    Account account2 =  
        new Account(4 );  
  
    System.out.println(  
        " Account with greater balance: "  
        + accountComparer .greater(account2 , account1 ));  
}
```

Java also provides a set of predefined functional interfaces for most common scenarios.

Questions

- What is Function Interface?
- What are the benefits of using Functional Interface?

LAMBDA EXPRESSION

Lambda expressions provide a convenient way to create an [anonymous class](#) . Lambda expressions implement [Functional Interface](#) more compactly. Lambda Expressions are primarily useful where you want to pass some functionality as an argument to another method and defer the execution of such functionality until an appropriate time.

Lambda expression can be just a block of a statement with method body and optional parameter types but without method name or return type. It can be passed as a method argument and can be stored in a variable.

```
// lambda expressions
```

```
() -> 123
```

```
(x,y) -> x + y
```

```
(Double x, Double y) -> x*y
```

Questions

- What is Lambda Expression?
- How are the Lambda Expression and Anonymous class-related?
- Can you pass Lambda Expression as a method parameter?
- What is the meaning of deferred execution of functionality, using a Lambda Expression?
- What are the benefits of using Lambda Expression?
- How are the Lambda Expression and Functional Interface related?

PURE FUNCTIONS

Pure functions are functions whose results depend only on the arguments passed onto them and are neither influenced by any change of state in the application nor altering the application state. Pure functions always return the same result for the same arguments.

```
public int increaseByFive( int original){  
    int toAdd = 5 ;  
    return original + toAdd ;  
}
```

Questions

- What is a Pure Function?
- What is the use of Pure Function in Functional Programming?
- How is the Pure Function guaranteed to always return the same results for the same arguments?

FLUENT INTERFACE

A **fluent interface** is used to transmit commands to subsequent calls, without the need to create intermediate objects and is implemented by method chaining. The **fluent interface** chain is terminated when a chained method returns **void**. **Fluent interface** improves readability by reducing the unnecessary objects created.

In the code below, the Fluent Interface is used to add a new Employee.

```
employee.create()  
    .atSite( "London" )  
    .inDepartment( "IT" )  
    .atPosition( "Engineer" );
```

Fluent interfaces are used primarily in scenarios where you construct queries, create a series of objects, or construct nodes in hierarchical order.

Questions

- What is the Fluent Interface?
- What are the benefits of defining and using Fluent Interface?
- Describe some usage of Fluent Interface?

GENERIC

GENERICIS

Generics is a mechanism that allows the same code to work on objects of different types in a type, class or interface, or a method; while providing compile-time protection. **Generics** are introduced to enforce the safety of type, in particular for collection classes.

- Once a parameter type is defined in **generics** , the object will work with the defined type only.
- In **generics** , the **formal type parameter** (E in the case below) is specified with the type (class or interface)

```
// Generic List type
// E is type parameter
public interface List < E > {
    void add( E x);
    Iterator < E > iterator();
}
```

- In **generics** , the **parameterized type** (**Integer** in this case) is specified when a variable of type is declared or the object of the type is created.

```
// variable of List type declared
// With Integer parameter type
List <Integer> integerList =
    new ArrayList<Integer>();
```

Notes

- There is a convention for naming parameters commonly used in **Generics** . E for element, T for type, K for the key, and V

- for values.
- A generic type is compiler support; all the type parameters are removed during the compilation and this process is called [erasure](#) .
- Due to strong type checking, generics help avoid many [ClassCastException](#) instances.

Generic Class Example

- The following NodeId generic class definition can be used with the object of any type.

```
// NodeId generic class with type parameter
public class NodeId< T > {
    private final T Id ;
    public NodeId( T id) {
        this . Id = id;
    }
    public T getId() {
        return Id ;
    }
}
```

- Usual SuperType – SubType rules do not apply to generics. In this example, even though [Integer](#) is derived from [Object](#) , still NodeId with type parameter [Integer](#) cannot be assigned to NodeId with type parameter [Object](#) .

```
// Parameter Type - Object
NodeId<Object> objectNodeId =
    new NodeId<>( new Object());

// Parameter Type - Integer
NodeId<Integer> integerNodeId =
    new NodeId<>( 1 );

// This results in an error
objectNodeId = integerNodeId ;
( error )
```

Questions

TYPE WILDCARDS

- A wildcard in generics is represented in the form of "?" . For example, a method which takes `List<?>` as a parameter will accept any type of List as an argument.

```
public void addVehicles( List <?> vehicles) {  
    //...  
}
```

- Optional **upper** and **lower** limits are set to enforce limitations, as the exact form of parameter defined by the wild card is not known.

```
// Only vehicles of type Truck can be added  
public void addVehicles  
    ( List <? extends Truck> vehicles) {  
    //...  
}
```

Notes

- Do not return wildcard in a return type as its always safer to know what is returned from a method.

Upper Bound

- To impose a restriction, upper bound can be set on the type parameters. Upper bound restricts a method to accept unknown type arguments extended only from the specified data type, like **Number** on the example below.

```
// Upper bound of type wild card  
public void addIds( List <? extends Number> T){
```

Lower Bound

- To impose a restriction, lower bound can be set on the type parameters. Lower Bound restricts a method to accept unknown type argument, which is a super-type of specified data type only, like Float in the example below.

```
// Lower bound of type wild card  
public void addIds( List <? super Float> T){
```

Type Inference

- Type inference is a compiler's ability to look at method invocation and declaration to infer the type arguments.
- In **Generics** , the operator called **Diamond operator** , <>, facilitates type inference.

```
// Type inference using Diamond operator  
List <Integer> integerList =  
    new ArrayList<>();
```

Questions

- What is Type Wildcard in Generics?
- Explain the Upper Bound type wildcard?
- Explain Lower Bound type wildcard?
- How are the different types of wildcards used in generics?
- What is automatic type inference in Generics? What is the diamond operator in Generics?

GENERIC METHOD

- Generic methods define their type parameters.

```
// generic method  
public < T > void addId ( T id){
```

- If we remove the `<T>` from the above method, we will get a compilation error as it represents the declaration of the type parameter in a generic method.
- The type (class or interface) that has a generic method does not have to be of generic type.
- While calling the generic method, we do not need to explicitly indicate the type parameter.

Notes

- Prefer using Generics and parameterized classes/methods to enforce compile-time type safety.
- Use the Bounded Type parameter to increase the flexibility of method arguments, at the same time it also helps to restrict the types that can be used.

Questions


- Explain the Generics method?
- Can you add a Generic method to a non-Generic type?

- What are the benefits of defining the bounded type as method parameters?
- How is compile type safety enforced by Generics?

JAVA GENERICS VS JAVA ARRAY

Java Generics

Consider the following hierarchy,



As both **Car** and **Bike** are derived from **Vehicle** , is it possible to assign `List<Car>` or `List<Bike>` to a variable of `List<Vehicle>` ?

No, `List<Car>` or `List<Bike>` cannot be replaced with `List<Vehicle>` , because you cannot put a **Bike** in the same list that has cars. So even though **Bike** is a **Vehicle** , it's not is a **Car** .

```
List< Bike > bikes =  
    new ArrayList< Bike >();
```

```
// suppose this is allowed  
List< Vehicle > vehicles = bikes ;  
vehicles .add( new Car ());
```

```
// Error - Bike and Vehicle are  
// considered incompatible types.
```

```
Bike bike = vehicles .get( 0 );
```

Java compiler does this checking compile time for incompatible types. For more on this refer to [Type Erasure](#) .

JAVA ARRAY

However, unlike Generics, the array of `Car` can be assigned to the array of `Vehicle` :

```
Car [] cars = new Car [ 3 ];  
Vehicle [] vehicles = cars ;
```

For Array, even though the compiler allows the above code to compile but when you run this code, you will get `ArrayStoreException` .

```
// this will result in ArrayStoreException  
vehicles [ 0 ] = new Bike ();
```

So even though the compiler did not catch this issue, the runtime type system caught it. An array is `reifiable` types, which means that run time is aware of its type.

Questions

- What are the reifiable types?
- Why can't you assign a Generic collection object of sub-type to a Generic collection object of super-type?
- Why it's allowed to assign an array object of sub-type to an array object of super-type? Are Java array more polymorphic?

TYPE ERASURE

For Java generic types, due to a process known as type [erasure](#) , Java compiler discards the type information and it is not available at run time. So because the type information is not available runtime, java compiler takes extra care to stop you at compile time itself, preventing any heap pollution.

Generics type is also of [non-reifiable](#) types, which means that its type information is removed during the compile time.

Questions

- What are non-reifiable types?
- What is type erasure?
- What would happen if type erasure is not there?

CO-VARIANCE

Covariance is a concept that you can read items from a generics defined with the upper bound type wildcard, but you cannot write anything into the collection.

Consider the following declarations with upper bound type wildcard:

```
List<? extends Vehicle> vehicles =  
    new ArrayList<Bike>();
```

You are allowed to read from `vehicles` generic collection because whatever is present in the list is sub-class of `Vehicle` and can be up-casted to a `Vehicle`.

```
Vehicle vehicle = vehicles.get(0);
```

However, you are not allowed to put anything into a covariant structure.

```
// This is a compile-time error  
vehicles.add(new Bike());
```

This would not be allowed, because Java cannot guarantee the actual type of the object in the generic structure. It can be anything that extends `Vehicle`, but the compiler cannot be sure. So you can read, but not write.

Questions

- Explain co-variance?
- Why can't you add an element of subtype to a generic defined with upper bound type wildcard?

CONTRA-VARIANCE

Contra-variance is the concept that you can write items to a generic defined with **lower bound type wildcard** , but you cannot read anything into the collection.

Consider the following declarations with lower bound type wildcard:

```
List<? super Car> cars = new ArrayList<Vehicle>();
```

In this case, even though the ArrayList is of type **Vehicle** but you can add **Car** into it through contra-variance; because **Car** is derived from **Vehicle** .

```
cars.add(new Car());
```

However, you cannot assume that you will get a **Car** object from this contra-variant structure.

```
// This is a compile-time error  
Car vehicle = cars.get(0);
```

Questions

- Explain contra-variance?
- Why can't you read an element from a Generic defined with lower bound type wildcard?

CO-VARIANCE VS CONTRA-VARIANCE

- Use `covariance` when you only intend to read generic values from the collection.
- Use `contra-variance` when you only intend to add generic values into the collection.
- Use the specific generic type when you intend to do both, read from, and write to the collection.

Questions

- When should you use co-variance?
- When should you use contra-variance?

COLLECTIONS

COLLECTIONS

Collections are the data structures that are basic building blocks to create any production level software application in Java. Interviewers are interested in understanding different design aspects related to the correct usage of collections. Each collection implementation is written and optimized for a specific type of requirement, and interview questions are to gauge the interviewee's understanding of such aspects.

Questions are often asked to check whether the candidate understands the correct usage of collection classes and is aware of alternative solutions available.

Following are few aspects on which questions on collections are asked:

- Collection types in Java.
- Unique features of different collection types.
- Synchronized collection.
- Concurrent collection.
- Ordering of elements in a collection.
- Speed of reading from collection.
- Speed of writing to a collection.
- The uniqueness of the elements in a collection.
- Ease of inter-collection operation.
- Read-only collections.
- Collection navigation.

COLLECTION FUNDAMENTALS

The collection is a container that groups multiple elements. Following is a simple example of a collection.

```
// Create a container list of cities
List <String> cities = new ArrayList<>();
// add names of cities
cities .add( "London" );
cities .add( "Edinburgh" );
cities .add( "Manchester" );
```

Notes

- Collections work with reference types.
- All collection interface implementation is [Generic](#) .
- Unlike arrays, all collection types can grow or shrink in size.
- Java provides a lot of methods to manipulate collections based on their use, so always test the existing methods before you implement one.

Collection Framework

The collection framework is defined by the following components.

- [Interfaces](#) - are the abstract types defined for each specific type of usage and collection type.
- [Implementation](#) - are concrete implementation classes to create an object to represent a different type of collection.

- **Algorithms** - are applied to the collections to perform various computations and to manipulate the elements in the collection.

Collection Framework helps you to reduce programming efforts, by providing data structures and **algorithms** to operate on them.

Questions

- Explain collections?
- How collections and Generics related?
- Can you use collection with primitive types?
- How can you use collection with the primitive types?
- Explain the difference between collections and arrays?
- What is the benefit of the collection framework?
- What are the different components of the collection framework?

COLLECTION INTERFACES

An interface defines its behavior in the form of the signature of the methods. To use a collection, you should always write code against collection interfaces and not the class implementations, so that the code is not tied to a specific implementation. This protects from possible changes in the underlying implementation class.

The following are the most important interfaces that define collections and their behavior. Each child node below is inherited from its parent node.

- + Collection
 - + Queue
 - + BlockingQueue
 - + TransferQueue
 - + BlockingDeque
 - + Deque
 - + BlockingDeque
 - + List
 - + Set
 - + SortedSet
 - + NavigableSet
- + Map
 - + SortedMap

Questions

- Why should you write code against the collection interface and not concrete implementation?

COLLECTION TYPES

- **Set** - Set contains unique elements.
- **List** - List is an ordered collection.
- **Queue** - Queue holds elements before processing in a FIFO manner.
- **Deque** - Deque holds elements before processing in both, FIFO and LIFO manner.
- **Map** - The map contains the mapping of keys to corresponding values.

Questions

- What are the different collection types?
- Define Set collection Type?
- Define List collection Type?
- Define Queue collection Type?
- Define Deque collection Type?
- Define Map collection Type?
- What is the difference between Queue and Deque?

Set

Basic Set

```
// Create a set
Set <String> set =
    new HashSet<>();
```

- **Set** is a collection of unique elements.
- Elements in the **Set** are stored unordered.
- Only one null element can be added to a **Set**.
- Duplicate elements are ignored.
- When ordering is not needed, **Set** is the fastest and has a smaller memory footprint.

Linked Hashset

```
// Create a linkedHashSet
Set <String> linkedHashSet =
    new LinkedHashSet<>();
```

- **LinkedHashSet** keeps the **Set** elements in the same order in which they were inserted.
- Insertion order is not affected in **LinkedHashSet** if an element is re-inserted.
- **Iterator** in **LinkedHashSet** returns elements in the same order in which these were added to the collection.

Sorted Set

- **SortedSet** imposes the ordering of elements to be either sorted in natural order by implementing a **Comparable** interface or custom sorted using a **Comparator** object.
- **TreeSet** is an implementation class for the **SortedSet** interface.

```
// Create a sorted set of city names
SortedSet <String> cityNames =
    new TreeSet<>();
```

- Use a [Comparator](#) object to perform custom sorting.

```
SortedSet <City> sortedCitiesByName =  
    new TreeSet<>( Comparator .comparing(  
        City::getName));
```

- If an element implements a [Comparable](#) interface, then the [compareTo\(\)](#) method is used to sort in the natural order.

[Navigable Set](#)

- [NavigableSet](#) inherits from the [SortedSet](#) and defines additional methods.
- [NavigableSet](#) can be traversed in both, ascending and descending order.
- [TreeSet](#) is one of the implementation classes for the [NavigableSet](#) interface.

[Questions](#)

- Can you add duplicate elements in a Set?
- Which collections type should you use when ordering is not a requirement? Why?
- Can you add a null element to a Set?
- Which Set class should you use to maintain the order of insertion?
- What happens to the order, if the same element is inserted again in a Set? Will it maintain its original position or inserted in the end?
- Which Set class should you use to ensure the ordering of the elements?
- Which Set class should you use when you need to traverse in both the directions?

List

- The `List` is an ordered collection of objects.
- The `List` can have a duplicate.
- The `List` can have multiple null elements.
- In `List`, an element can be added at any position.
- Using `ListIterator`, a list can be iterated in both, forward and backward direction.

```
// Create a list of cities  
List <String> cities =  
    new ArrayList<>();
```

ArrayList

- `ArrayList` is based on an array.
- It performs better if you access elements frequently.
- Add and remove is slower in the `ArrayList`. If an element is added anywhere but the end requires shifting of the element. If elements are added beyond its capacity then the complete array is copied to a newly allocated place.

Linked List

- `LinkedList` is based on the list.
- `LinkedList` is slower in accessing elements and only sequential access allowed.
- Adding and removing elements from `LinkedList` is faster.
- `LinkedList` consumes more memory than `ArrayList` as it keeps pointers to its neighboring elements.

List Iterator

- `ListIterator` iterates list in both the directions

```
// full list iterator
ListIterator <String> fullIterator =
    cityList .listIterator();

// partial list iterator starts at index 3
ListIterator <String> partialIterator =
    cityList .listIterator( 3 );
```

Notes

- Random access is better in [ArrayList](#) as it maintains an index-based system for its elements whereas [LinkedList](#) has more overhead as it requires traversal through all elements.

Questions

- Can you add duplicate elements to a List?
- Can you add null to a list?
- In which scenarios would you use ArrayList in: frequent add/update or frequent reading?
- In which scenarios would you use LinkedList in: frequent add/update or frequent reading?
- Which one takes up more memory between ArrayList and LinkedList? Why?
- Between ArrayList and LinkedList, which one would you prefer if you need random access? Why?

Queue

- A queue is a collection designed to hold elements before processing.
- A queue has two ends, a tail, and a head.
- A queue works in a FIFO manner, first in and first out.

Types of Queues

- Queue - simple queue which allows insertion at tail and removal from the head, in a LIFO manner.
- Deque - allows insertion and removal of elements from head and tail.
- Blocking Queue - blocks the thread to add element when its full and also blocks the thread to remove element when it's empty.
- Transfer Queue - special blocking queue where data transfer happens between two threads.
- Blocking Deque - a combination of Deque and blocking queue .
- Priority queue - element with the highest priority is removed first.
- Delay queue - element is allowed to be removed only after the delay associated with it has elapsed.

Basic Queue

- A queue has one entry point (tail) and one exit point (head).
- If entry and exit point is the same, it's a LIFO (last in first out) queue .

```
// simple queue
Queue <String> queue =
    new LinkedList<>();
```

Priority Queue

- In the PriorityQueue, a priority is associated with the elements in the queue.

- Element with the highest priority is removed next.
- `PriorityQueue` does not allow the `null` element.
- An element of queue either implement a `Comparable` interface or use a `Comparator` object to calculate priority.

```
// City class implements Comparable interface
Queue <City> pq =
    new PriorityQueue<>();
```

Deque

- Deque allows insertion and removal from both ends.
- `Deque` does not provide indexed access to elements.
- `Deque` extends the `Queue` interface.

```
// Create a Deque
Deque <String> deque =
    new LinkedList<>();
```

Blocking Queue

- `BlockingQueue` interface inherits from `Queue` .
- `BlockingQueue` is designed to be thread-safe.
- `BlockingQueue` is designed to be used as producer-consumer queues.

Transfer Queue

- `TransferQueue` extends `BlockingQueue` .
- It may be capacity bounded, where `Producer` may wait for space availability and/or `Consumers` may wait for items to become available.

```
// Transfer Queue
TransferQueue <String> ltq =
    new LinkedTransferQueue<String>();
```

Questions

- What is the purpose of the Queue?
- How is Deque different from Queue?
- What is a priority queue?
- What criteria should be used to delete an item from the priority queue?
- What is a Delay Queue?
- What is a Blocking Queue?
- Which Queue is a thread-safe queue?
- Which Queue implements the Producer-Consumer pattern?
- Explain the difference between the Blocking Queue and the Transfer Queue?
- Which Queue is capacity bounded, allowing only the specified number of elements in Queue? What happens if you try to put more than the capacity?

Map

- The `map` contains key-value mapping.
- Usually, `Map` allows one null as its key and multiple null as values, but it's left to the Map's implementation class to define restrictions.

Implementation

HashMap

- `HashMap` is based on a hash table.
- `HashMap` does not guarantee the order of the elements in the map.
- `HashMap`'s hash function provides constant-time performance for the `get()` and the `put()` operations.
- `HashMap` allows one `null` for the key and multiple `nulls` for the value.

```
// Create a map using HashMap
Map <String, String> hashMap =
    new HashMap<>();
```

LinkedHashMap

- `LinkedHashMap` is a hash table and linked list implementation of `Map` interface.
- `LinkedHashMap` stores entry using a doubly-linked list.
- Use `LinkedHashMap` instead of `HashSet` if the insertion order is to be maintained.
- The performance of `HashMap` is slightly better than `LinkedHashMap` as `LinkedHashMap` has to maintain the linked list too.
- It ensures iteration over entries in its insertion order.

```
// LinkedHashMap
LinkedHashMap lhm =
    new LinkedHashMap();
```

WeakHashMap

- **WeakHashMap** stores only weak references to its keys.
- **WeakHashMap** supports both **null** key and **null** values.
- When there is no reference to keys, they become a candidate for garbage collection.

```
// WeakHasMap  
Map map =  
    new WeakHashMap();
```

Sorted Map

- **SortedMap** provides complete ordering on its keys.
- Sorts the map entries on keys based either on natural sort order (**Comparable**) or custom sort order (**Comparator**).
- **SortedMap** interface inherits from **Map**.

```
// sorted map  
SortedMap <String,String> sm =  
    new TreeMap<>();
```

Navigable Map

- **NavigableMap** extends the **SortedMap** interface.
- **TreeMap** class is the implementation of **NavigableMap**.
- **NavigableMap** can be accessed in both, ascending and descending order.
- It supports navigation in both directions and getting the closest match for the key.

```
// Create a Navigable Map  
NavigableMap <String,String> nm =  
    new TreeMap<>();
```

Concurrent Map

- `ConcurrentHashMap` is concrete implementation of `ConcurrentMap` interface.
- `ConcurrentMap` uses a fine-grained synchronization mechanism by partitioning the map into multiple buckets and locking each bucket separately.
- `ConcurrentHashMap` does not lock the map while reading from it.

```
// Create a Concurrent Map
ConcurrentMap <String,String> cm =
    new ConcurrentHashMap<>();
```

Notes

- The `Map` keys and `Set` items must not change state, so these must be immutable objects.
- To avoid implementation of `hashCode()` and `equals()`, prefer using immutable classes provided by JDK as key in `Map`.
- Never expose collection fields to the caller, instead provide methods to operate on those.
- `HashMap` offers better performance for inserting, locating, and deleting elements in a map.
- `TreeMap` is better if you need to traverse the keys in sorted order.

Questions

- How Map is different from other collections?
- Can you add a null value as the key to a Map?
- Can you guarantee ordering of elements in HashMap?
- How HashMap can provide constant-time performance for the `get()` and the `put()` operations.
- What is the difference between HashMap and LinkedHashMap?
- If the insertion order is to be maintained, which one should you use: LinkedHashMap or HashSet? Why?

- What is WeakHashMap?
- If do you want to maintain the ordering of key, which Map class should you use?
- Which Map class should you use to access it in both ascending and descending order?
- Which Map implement is optimized for thread safety?
- How does ConcurrentHashMap enforce scalability, with good efficiency still maintained?
- Why the key used for Map and value used for Set should be of immutable type?
- What are the disadvantages of exposing an internal collection object to the caller?
- Which collection is best for traversing in sorted order?

ALGORITHMS

- **Sorting** - is used to sort the collection in ascending or descending order. You can either use the natural order if the key class has implemented a **Comparable** interface or need to pass the **Comparator** object for custom sorting.
- **Searching** - **BinarySearch** algorithm is used to search keys.
- **Shuffling** - it re-orders the elements in the list in random order; it is generally used in games.
- **Data Manipulation** - reverse, copy, swap, fill, and addAll algorithms are provided to manipulate the usual data on the collection elements.
- **Extreme Value** - **min** and **max** algorithms are provided to find the minimum and maximum values in the specified collection.

Questions

- What are the different algorithms supported by the collection framework?
- Which search algorithm is used to search keys?
- How is the shuffling algorithm used?

COMPARABLE VS COMPARATOR

`Comparable` interface is implemented to sort the collections in a natural order and the `Comparator` object is used to perform a custom sort on the collections.

If an element implements a `Comparable` interface, then the `compareTo()` method is used to sort in the natural order. The `compareTo()` method compares the specified object with the existing object for order. It returns a negative integer, zero, or a positive integer as the existing object is smaller than, equal to, and greater than the specified object.

Comparable interface implementation

```
@Override
public int compareTo( Employee employee) {
    //comparison logic
}
```

A `Comparator` object contains logic to compare two objects that might not align with the natural ordering. `Comparator` interface has `compare()` method, which takes two objects to return a negative integer, zero or a positive integer as the first argument is smaller than, equal to and greater than the second.

Using a `Comparator` object to perform custom sorting.

```
SortedSet <City> sortedCitiesByName =
    new TreeSet<>( Comparator .comparing(
        City::getName));
```

Comparator is generally used to provide a custom comparison algorithm in a situation when you do not have complete control over the class.

Questions

- What is the difference between the Comparable Interface and Comparator Class?
- When should you implement Comparable Interface to sort a collection?
- When do you use Comparator Class to sort a collection?

HASHCODE() AND EQUALS()

- The `HashTable` , `HashMap`, and `HashSet` uses `hashCode()` and `equals()` methods to calculate and compare objects that are used as their key.
- All classes in Java inherit the base hash scheme from the `Object` base class unless they override the `hashCode()` method.
- Any class that overrides `hashCode()` method is supposed to override `equals()` method too.

Implementation of hashCode() and equals()

```
final public class Employee {  
  
    final private int id ;  
    final private String name ;  
    final private String department ;  
  
    @Override  
    public boolean equals( Object o) {  
        if ( this == o) return true ;  
        if (o == null || getClass() !=  
            o.getClass())  
            return false ;  
  
        Employee employee = ( Employee ) o;  
  
        if ( id != employee . id ) return false ;  
        if ( name != null ?  
            ! name .equals( employee . name ) :
```

```

        employee . name != null )
return false ;

return !( department != null ?
! department .equals( employee . department ) :
employee . department != null );

}

@Override
public int hashCode() {
    int result = id ;
    result = 31 * result +
        ( name != null ?
            name .hashCode() :
            0 );
    result = 31 * result +
        ( department != null ?
            department .hashCode() :
            0 );
    return result ;
}
}

```

Rules of hashCode and equals

- Once an object is created, it must always report the same `hashCode()` during its lifetime.
- Two equal objects of the same class shall return the same hash-code.


Note

Two objects, which are not equal, can have the same hash-code.

How hashCode() and equals() are used

- The `hashCode()` and `equals()` method is used to segregate items into separate buckets for lookup efficiency. If two items have the same hash-code then both of these will be stored in the same bucket, connected by a linked list.
- If the hash-code of two objects is different, then the `equals()` method is not called; otherwise, the `equals()` method is called to compare the equality.

The figure below conceptually demonstrates the storage of objects in a hash bucket.



Questions

- What are the different collection types in Java that use hashCode and equals methods? Why do these collections type use hashCode and equals method?
- If you override the hashCode method, why should you always override the equals method too?
- The hashCode method is already present in the java.lang.Object class, why should you override it?
- Why does an object have to return the same hash-code value every time during its lifetime?
- Is it necessary that an object returns the same hash-code every time program runs?
- Is it necessary that two equal objects return the same hash-code?
- Can two objects, that are not equal, still have the same hash-code?
- What is the purpose of calculating hash-code?
- What is the purpose of the equals method?
- If two objects have the same hash code, how are these stored in hash buckets?
- If two objects have different hash codes, do you still need to implement an equals method?
- If there are two objects with the same hash code, then how the equals method is used?

HASHTABLE VS HASHMAP

[HashTable](#) and [HashMap](#) are data structures used to keep the key-value pair. These maintain an array of buckets and each element is added to a bucket based on the hashcode of the key object.

The major difference between these is that the [HashMap](#) is non-synchronized. This makes HashMap better for single-threaded applications, as unsynchronised Objects perform better than synchronized ones due to lack of locking overhead. In a multi-threaded application, HashTable should be used.

A HashMap can also be converted to a synchronized collection using the following method:

```
Collections.synchronizedMap( idToNameMao );
```

Questions

- What is the purpose of HashTable and HashMap?
- What is the difference between HashTable and HashMap?
- How do HashTable and HashMap store objects?
- Why HashMap is better for single-threaded applications?
- Why performance of HashMap better than HashTable?
- How can you use HashMap in a multi-threaded environment?

SYNCHRONIZED VS CONCURRENT COLLECTIONS

Both [synchronized](#) and [concurrent](#) collection classes provide thread safety, but the primary difference between these is in terms of scalability and performance.

Collection Synchronization

- [Collection synchronization](#) is achieved by allowing access to the collection only from one thread, at a time.
- Synchronization is achieved by making all the public methods [synchronized](#) .
- Any composite operation that invokes multiple methods, needs to be handled with locking the operation at the client-side.
- Examples of synchronized collections are [HashTable](#) , [HashSet](#), and [synchronized HashMap](#) .

Collection Concurrency

- [Collection concurrency](#) is achieved by allowing simultaneous access to the collection from multiple threads.
- [ConcurrentHashMap](#) implements very fine-grained locking on collection by partitioning the collection into multiple buckets based on hash-code and using different locks to guard each hash bucket.
- The performance is significantly better because it allows multiple concurrent writers to the collection at a time, without locking the entire collection.

- Examples of concurrent collection are `ConcurrentHashMap` , `CopyOnWriteArrayList` and `CopyOnWriteHashSet` .

For large collections prefer using a concurrent collection like `ConcurrentHashMap` instead of `HashTable` as the performance of `Concurrent` collection will be better due to less locking overhead.

Note that in multi-threaded scenarios, where there are interoperation dependencies, you still need to provide synchronized access to these composite operations on concurrent collections, as depicted below:

```
Object synchObject = new Object();

ConcurrentHashMap<String, Account> map =
    new ConcurrentHashMap<>();

public void updateAccount( String userId){
    synchronized ( synchObject ) {
        Account userAccount = map .get(userId);
        if ( userAccount != null ) {
            // some operation
        }
    }
}
```

Questions

- What is the purpose of Synchronized and Concurrent collections?
- What is collection synchronization?
- What is collection concurrency?
- How do you achieve collection synchronization?
- How do you ensure coordinated collection access when a composite operation needs several methods to complete the task?
- How do you achieve collection concurrency?

- How does ConcurrentHashMap offer multiple threads concurrent access to the same collection?
- Why performance of ConcurrentHashMap is better even though it allows simultaneous writes from multiple threads?
- What would you prefer between the ConcurrentHashMap and the HashTable? Why?

ITERATING OVER COLLECTIONS

There are several ways to iterate over a collection in Java. The following are the most common methods.

Iterable.forEach

The `forEach` method is available with the collections implementing the `Iterable` interface, and the `forEach` method takes a parameter of type `Consumer` .

```
List< Account > accounts =  
    Arrays.asList(  
        new Account( 123 ),  
        new Account( 456 ));  
  
accounts.forEach(acc -> print(acc));
```

for-each looping

- A for-each loop can be used to loop the `map.entrySet()`, to get key and value both.

```
for (Map.Entry< String , Account > accountEntry :  
    map.entrySet()) {  
    print( "UserId - " + accountEntry.getKey() + ", " +  
        "Account - " + accountEntry.getValue());  
}
```

- A for-each loop can be used to loop the `map.keySet()` to get keys.

- A for-each loop can be used to loop `map.values()` to get values.
- While running a `for-each` loop, collection cannot be modified.
- A `for-each` loop can only be used to navigate forward.

Iterator

- The `Iterator` can move in both directions backwards and forward.
- When using the `Iterator`, you can remove entries during an iteration, which is not possible when you use a `for-each` loop.
- The `for-each` loop also uses `Iterator` internally.

```
Iterator<Map.Entry< String , Account >> accountIterator =
    map .entrySet().iterator();

while ( accountIterator .hasNext()){
    Map.Entry< String , Account > accountEntry =
        accountIterator .next();

    print( "UserId - " + accountEntry .getKey() + ", " +
        "Account - " + accountEntry .getValue());
}
```

Notes

- The `for-each` loop should be preferred over the `for` loop, as the `for` loops can be the source of errors, specifically related to index calculations.
- The `Iterator` is considered more thread-safe because it throws an exception if the collection changes during iteration.

Data Independent Access

The code should be written in such a way that client code should not be aware of the internal structure used to store the collection. This enables making internal changes without breaking any client code. So to facilitate

data-independent access to the collection, it must be exposed such that the **Iterator** can be used to iterate through all the elements of the collection.

Questions

- Explain different ways to iterate over collection?
- Can you modify the collection structure while iterating using a for-each loop?
- Can you modify a collection element value while iterating using a for-each loop?
- What are the limitations with the for-each loop, regarding the navigation direction?
- What is different between the for-each loop and the Iterator, regarding the navigation direction?
- Can you modify the collection structure while iterating using an Iterator?
- Can you modify a collection element value while iterating using an Iterator?
- Between the for-each loop and the for loop, which one should you prefer? Why?
- Why Iterators are considered more thread-safe?
- Explain the concept of Data Independent Access of the collection?
- How can you design your class to provide Data Independent Access to the collections used in the class?

FAIL FAST

The [Iterator](#) methods are considered [fail-fast](#) because Iterator guards against any structural modification made to the. This ensures that any failure is reported quickly rather than application landing into a corrupt state sometime later. In such a scenario, [ConcurrentModificationException](#) is thrown. The Iterators from [java.util](#) are fail fast.


Questions

- What is the term fail-fast used in collection iteration?
- What is the benefit of a fail-fast iterator?
- What is a fail-safe iterator?
 - *The fail-safe iterator doesn't throw any exception when a collection is modified, because fail-safe iterator works on the clone of the original collection. Iterators from [java.util.concurrent](#) package is fail-safe.*

ERROR AND EXCEPTION HANDLING

EXCEPTION

Exception class hierarchy in java



- An **exception** is an abnormal situation that interrupts the flow of program execution.
- All exceptions inherit from **Throwable**.
- You subclass **Exception** class if you want to create a **checked** exception or **RuntimeException** if you want to create an **unchecked** exception.
- Though you can theoretically subclass **Throwable** class to create **checked** exception, this is not recommended as **Throwable** is super-class for all Java exceptions and errors.

Questions

- What is an Exception?
- What are the root level exception super classes in Java?
- Which super-class should you sub-class to create checked exception?
- Which super-class should you sub-class to create unchecked exceptions?

CHECKED VS UNCHECKED VS ERROR

Checked Exception

- The checked exceptions are checked at compile time.
- Except for `RuntimeException` and `Error` classes, the checked exceptions extend `Throwable` or its sub-classes.
- The checked exceptions are programmatically recoverable.
- You can handle checked exceptions either by using `try/catch` block or using the `throws` clause in the method declaration.
- Static initializers cannot throw checked exceptions.

```
public static void main(String args[]) {  
    FileInputStream fis = null ;  
    try {  
        fis = new FileInputStream("details.txt" );  
    } catch (FileNotFoundException fnfe) {  
        System.out.println("Missing File :" + fnfe);  
    }  
}
```

Unchecked Exception

- The unchecked exceptions are checked at runtime.
- Unchecked exceptions extend `RuntimeException` class.
- You can't fairly hope to recover from those exceptions.
- The unchecked exceptions can be avoided using good programming techniques.
- Throwing unchecked exceptions helps to reveal other shortcomings.

```

public int getAccountBalance(
    String customerName) {
    int balance = 0 ;
    if (customerName == null )
        throw new IllegalArgumentException("Null argument" );
    // logic to return calculate balance
    return balance ;
}

```

Error

- The **error** classes are used to define irrecoverable and fatal exceptions, which the applications are not supposed to catch.
- Generally programmers cannot do anything to recover from these exceptions.
- Even if you catch **OutOfMemoryError** , you may get it again because there is a high probability that the Garbage Collector may not be able to free memory.

Use **checked** exception only for the scenario where failure is expected and there is a very reasonable way to recover from it; for anything else use **unchecked** exception.

Questions

- What are the checked exceptions?
- What are unchecked exceptions?
- What types of exceptions does the Error class in Java defines?
- How can you handle checked exceptions?
- What happens if an exception is unhandled?
- What are the different ways to handle checked exceptions?
- Which exception classes will you use in the catch block to handle both: checked and unchecked exceptions?
- How do you choose between checked and unchecked exceptions?
- Can you recover from unchecked exceptions?
- How can you avoid unchecked exceptions?
- Can you throw checked exceptions from the static initializer block? Why?

- *You cannot throw because there is no specific place to catch it and it's called only once. You have to use try/catch to handle checked exceptions .*
- What should you do to handle an Error?

EXCEPTION HANDLING

BEST PRACTICES

Even though exception handling is primarily driven by context, but there must be a consistency in the exception handling strategy. Following are few exception handling best practices:

1. [Never suppress an exception](#) - as it can lead your program to an unsafe and unstable state.
2. [Don't perform excessive exception handling](#) - specifically when you do not know how to completely recover from it.
3. [Never swallow an exception](#) - as it may lead the application into an inconsistent state, and even worst, without recording the reason for it.
4. [Don't catch and continue program execution](#) - with some default behavior. Default behavior defined today may not be valid in the future.
5. [Don't show a generic error message to the user](#) - instead, clean the exception handling code to report user-friendly message with a suggestion about the next step.
6. [Don't put more than one exception scenarios in single try-catch](#) - as it will be impossible to ascertain the reason for the exception.
7. [Don't catch multiple checked exceptions in a single catch block](#) - as it will be impossible to ascertain the reason for the exception.
8. [Don't unnecessarily wrap the exception](#) - which may mask the true source.
9. [Don't reveal sensitive information](#) - instead sanitize exceptions generated specifically from the sources that may reveal sensitive information.

10. [Always log exception](#) - unless there is a compelling reason not to do so.
11. [Don't catch Throwable](#) - as it will be impossible to ascertain the reason for the exception.
12. [Don't use exception to control the flow of execution](#) - instead, use a boolean to validate a condition where possible.
13. [Handle different scenarios programmatically](#) - instead of putting all coding logic in a try block.
14. [Explicitly name the threads](#) - in a multithreaded application, it significantly eases the debugging.
15. [Never throw a generic exception](#) - as it will be impossible to ascertain the reason for the exception.

Questions

- Describe some exception handling best practices?
- What are the pitfalls of suppressing an exception?
- What is the problem with showing a generic error message?
- What is the downside of swallowing an exception?
- What are the pitfalls of handling multiple exceptions in a single catch block?
- What would you do if an exception is thrown from a source that contains sensitive information? What would you log in such a case and what message would you show to the user?
- What is the pitfall of wrapping all the exceptions into a Generic exception class?
- Why shouldn't you use exceptions to control the flow of program execution?
- Should you log all the exceptions? Why?

- Why shouldn't you use Throwable or some other root level class to catch exceptions?
- What criteria should be used to decide which block of code should be included in the try block?
- If a nested call is made that goes through multiple methods, would you be using a try-catch for each method? Why?

TRY-WITH-RESOURCE

`try-with-resource` is a Java language construct, which makes it easier to automatically close the resources enclosed within the `try` statement.

```
try (FileInputStream fis =  
    new FileInputStream("details.txt" )) {  
    // Code to read data  
}
```

- The resource used with `try-with-resources` must inherit the `AutoCloseable` Interface.
- You can specify multiple resources within a try block.

Questions

- Which java construct can you use to close the system resources automatically?
- Which Interface resources must inherit to use a class object within a try-with-resource construct?

THREADING

THREADING TERMS

- A **thread** is the smallest piece of executable code within a process.
- A **program** is a set of ordered operations.
- A **process** is an instance of a program.
- **Context Switch** is an expensive process of storing and restoring the state of a thread.
- **Parallel processing** is the simultaneous execution of the same task on multiple cores.
- **Multithreading** is the ability of a CPU to execute multiple processes or threads concurrently.
- **Deadlock** occurs when two threads are waiting for each other to release the lock.

Basic Concepts

- All Java programs begin with the **main()** method on a user thread.
- The program terminates when there is no user thread left to execute.
- Thread maintains a private stack and series of instructions to execute.
- Thread has a private memory called thread-local storage, which can be used to store thread's current operation-related data, in a multi-threaded environment.
- JVM allows the process to have multiple threads.
- Each thread has a priority.

Questions

- What is a thread?

- What is a program?
- What is a process?
- What is the difference between a thread and a process?
- What is the difference between program and process?
- Explain context switching of thread?
- What is parallel processing?
- What is multi-threading?
- How parallel processing and multithreading related?
- What is deadlock?
- What is a user thread?
- What is thread-local storage? What type of objects should you would store with thread local storage?
- Do threads share stack memory?

THREAD LIFECYCLE

The following are various stages of thread states.

New - Thread is created but not started.

Runnable - Thread is running.

Blocked - Thread waiting to enter the critical section.

Waiting - Thread is waiting by calling `wait()` or `join()`.

Time-waiting - Thread waiting by calling `wait()` or `join()` with a specified timeout.

Terminated - Thread has completed its task and exited.

Notes

You can get the state of thread using the `getState()` method on the thread.

Questions

- Describe the different stages of the thread lifecycle?
- What is the difference between the blocked state and the waiting state?
- How can you find a thread state?
- How thread sleep is different from thread wait?

THREAD TERMINATION

The thread should be stopped calling `interrupt()` method. Calling interrupt on a thread even breaks it out of `Thread.sleep()` state.

The long operation executing on the thread should recurrently call `isInterrupted()` method to check whether the thread is requested to be stopped, where you can safely terminate the current operation and perform any required cleanup.

```
if (Thread.currentThread().isInterrupted()) {  
    // cleanup and stop execution  
    //, for example, a break in a loop  
}
```

Notes

The `stop()`, `suspend()`, and `resume()` methods are deprecated, as using these may lead the program to an inconsistent state.

Questions

- Define a good strategy to terminate a thread?
- What is thread interrupt? How thread's interrupt method is different from the thread stop method?
- What happens if a thread is sleeping and you call interrupt on the thread?
- Does calling interrupt stop the thread immediately?

IMPLEMENTING RUNNABLE VS EXTENDING THREAD

Thread instantiated implementing `Runnable` Interface

```
public class RunnableDerived  
    implements Runnable {  
    public void run() {  
    }  
}
```

Thread created extending `Thread` class

```
public class ThreadDerived  
    extends Thread {  
    public void run() {  
    }  
}
```

- `Runnable's run()` method does not construct a new thread but runs as a normal process in the same thread on which it was created, whereas the `Thread's start()` methods construct a new thread.
- `Runnable` is the preferred way to create a thread unless you are specializing `Thread` class, which is unlikely in most of the case.
- By Implementing `Runnable` you are providing a specialized class an additional ability to run too.
- Also by separating the task as `Runnable` you can execute the task using different means.

Questions

- What are the different ways to create a thread?
- How implementing the Runnable interface is different from extending Thread class?
- When should you extend the Thread class?
- When should you inherit Runnable Interface?
- Which one is the preferred way between Runnable and Thread?
- Does implement Runnable creates a thread?
- Both Thread class and Runnable Interface have run methods, what is the difference?

RUNNABLE VS CALLABLE INTERFACE

Runnable Interface

```
@FunctionalInterface
public interface Runnable {
    public abstract void run();
}
```

Callable Interface

```
@FunctionalInterface
public interface Callable<V> {
    V call() throws Exception ;
}
```

- A **Runnable** cannot return a result and cannot throw a checked exception.
- A **Callable** needs to implement a **call()** method while a **Runnable** needs to implement the **run()** method.
- A **Callable** can be used with **ExecutorService** methods but a **Runnable** cannot be.

Questions

- What is the difference between the Callable and Runnable interface?
- What is the benefit of using Runnable over Callable?
- Can you throw checked exceptions from the Runnable interface?

- Why Runnable and Callable interfaces called Functional Interface?
- Which Interface returns the result: Runnable or Callable?

DAEMON THREAD

A **daemon thread** is a thread, which allows **JVM** to exit as soon as the program. As soon as JVM halts, all the daemon threads exist without unwinding stack or letting any **finally** blocks to execute. Daemon thread executes on very low priority.

Threads inherit the daemon status or parent thread, which means that any thread that is created by the main thread will be a non-daemon thread.

Generally, the daemon threads are used to support background tasks or services for the application. Garbage Collection happens on Daemon thread.

You can create a daemon thread like the following:

```
Thread thread = new Thread();  
thread.setDaemon(true);
```

All non-daemon threads are called user threads. Those threads stop the JVM from closing.

The process terminates when there are no more user threads. The Java main thread is always a user thread.

Questions

- What is a daemon thread?
- What is user thread? What is the main thread?
- Can a program exit if daemon thread is still running?
- Can a program exit if the user thread is still running?

- Does a finally block on a Daemon thread still get a chance to execute When THE JVM halts exiting all the Daemon thread?
- When a new thread is created; is it created as a user thread or a Daemon thread?
- What happens if no user thread is running but a Daemon thread is still running?
- What is Java's main thread: a user thread or a Daemon thread?
- What type of thread is created when you create a new thread on a main thread: daemon or user?
- What kind of work you can use Daemon thread for?

RACE CONDITION AND IMMUTABLE OBJECT

A [race condition](#) arises when multiple threads access shared data simultaneously to modify it. As the order of data read and write by these threads cannot be predicted, this can lead to an unpredictable data value.

An object is considered [immutable](#) when there is no possibility of its state change after the construction. If an object is immutable, it can be shared across multiple threads without worrying about the [race conditions](#).

To make an object Immutable

- Declare the class [final](#).
- Allow only constructor to create an object. Don't provide field [setter](#).
- Mark all the field [private](#).

Questions

- How immutable objects help to prevent race conditions?
- Why is race condition capable of generating unpredictable results?
- Why immutable objects are considered safe in a multi-threaded environment?
- Why it's suggested to declare an immutable class as final?
- Why constructor should be the only way to construct an immutable object? What happens if setters are provided?

THREAD POOL

[Thread Pool](#) is a set of a number of worker threads that exist separately from the [Runnable](#) and [Callable](#) tasks.

A fixed thread Pool reduces the overhead of thread creation. It helps the application to degrade gracefully when there is a surge in request that go beyond its capacity to process, by preventing an application from entering into a hanging state or crashing.

Thread Pool also enables a loosely coupled design by decoupling the creation and execution of tasks.

Creating a fixed thread pool is easy with the help of [Executors](#) class, where you can use the [newFixedThreadPool\(\)](#) factory method to create [ExecutorService](#) to execute tasks.

Questions

- What is a thread pool?
- How a thread pool reduces the overhead of thread creation?
- How does a thread pool help to prevent an application from hanging or crashing?
- What is a fixed thread pool and how is it created?
- How does the thread pool enables loosely coupled design?

SYNCHRONIZATION

CONCURRENT VS PARALLEL VS ASYNCHRONOUS

Parallel processing is the simultaneous execution of the same task on multiple cores.

Concurrent processing is the simultaneous execution of multiple tasks; either on multiple cores or by a pre-emptively time-shared thread on the processor.

Asynchronous processing is independent execution of a process, without waiting for a return value from intermediate operations.

Questions

- Explain concurrent processing?
- Explain parallel processing?
- Explain asynchronous processing?
- Explain the difference between concurrent and parallel processing?
- Does parallel processing require multiple threads?
- When an application is concurrent but not parallel?
 - *When application processes multiple operations simultaneously without dividing these operations further into smaller tasks.*

- When an application is parallel but not concurrent?
 - *When application processes one operation dividing the operation into smaller tasks that are processed in parallel.*
- When an application is neither concurrent, not parallel?
 - *When application processes only one operation without dividing the operation into smaller tasks.*
- When an application is both concurrent and parallel?
 - *When application processes multiple operations simultaneously and also dividing these operations further into smaller tasks that are processed in parallel.*

THREAD SYNCHRONIZATION

A **race condition** occurs when multiple threads concurrently access shared data to modify it. As it is not possible to predict the order of data read and write by these threads, it may lead to an unpredictable data value.

Critical Section is the block of code that if accessed concurrently, by more than one thread, may have undesirable effects on the outcome.

Thread Synchronization is controlling access to the **critical sections** to prevent undesirable effects in the program.

Synchronization establishes a memory buffer, known as **happen-before**, which ensures that any other thread that subsequently acquires the same local objects may access all changes made by the thread to the local objects in the critical sections.

Questions

- Explain race condition in multi-threading?
- What is a critical section?
- What is thread synchronization?
- What is a memory barrier?
- What is the concept of happen-before in thread synchronization?

SYNCHRONIZED METHOD VS SYNCHRONIZED BLOCK

The `synchronized` keyword is used to mark a critical section in the code. Mutual exclusion synchronization is achieved using locking the critical section using the `synchronized` keyword.

This can be done in the following two ways.

Marking a method as a critical section

```
public class DatabaseWrapper {  
    Object reference = new Object();  
    // Method marked as critical section  
    public synchronized void writeX() {  
        // code goes here  
    }  
    // Method marked as critical section  
    public static synchronized void writeY() {  
        // code goes here  
    }  
}
```

Marking block of code as a critical section

```
public void writeToDatabase() {  
    // multiple threads can reach here  
    // Code marked a critical section  
    synchronized (this) {  
        // only one thread can  
        // execute here at a time  
    }  
    // multiple threads can execute here  
}
```

Notes

- Minimize the scope of locking to just a **critical section** . This will improve overall performance and minimize chances for encountering a race condition.
- Prefer **synchronized** block over **synchronized** method, as block locks only on a local object as opposed to an entire class object.
- From within **synchronized** , never call a method provided by the client code or the one that is designed for inheritance.

Questions

- How is the synchronized keyword used?
- What is the difference between synchronized methods and synchronized block?
- Why synchronized block is preferred to synchronized method?
- Does the synchronized method lock the entire object?
- What happens if you call a method supported by the client from within a synchronized block or process?

CONDITIONAL SYNCHRONIZATION

Conditional [synchronization](#) is achieved using conditional variable along with [wait\(\)](#) and [notify\(\)](#) or [notifyAll\(\)](#) methods.

```
// conditional synchronization
public void operation()
    throws InterruptedException {
    synchronized (reference) {
        if (condition1) {
            // wait for notification
            reference.wait();
        }
        if (condition2) {
            // Notify all waiting threads
            reference.notifyAll();
        }
    }
}
```

1. There are two methods to signal the waiting thread(s).
 - a. [notify\(\)](#) - signals only one random thread.
 - b. [notifyAll\(\)](#) - signals all threads in a wait state.
2. The [wait\(\)](#) method has an overload to pass timeout duration too, [wait\(long timeout\)](#) .
3. Between [notify\(\)](#) and [notifyAll\(\)](#) method, prefer using [notifyAll\(\)](#) as it notifies all the waiting threads.
4. The [notify\(\)](#) method wakes a single thread, and if multiple threads are waiting to be notified, then the choice of thread is arbitrary.

Questions

- What is conditional synchronization?
- What is the propose of the wait call?
- What is the difference between notify and notifyAll method?
- With multiple threads waiting for the notification which one will be notified when you call to notify?
- Which one should you prefer between notify and notifyAll ?
Why?

VOLATILE

In a multi-threaded application, every thread maintains a copy of the variable from the main memory to its CPU cache. So any change made by a thread to the variable in its CPU cache will not be visible to other threads.


A field marked `volatile` is stored and read directly from the main memory. As `volatile` fields are stored in main memory, all the threads have visibility to the most updated copy of the `volatile` field's value, irrespective of the thread modified it.

Consider a class `Ledger`, which has a member `currentIndex` to keep track of the number of entries made. In a multi-threaded environment, each thread will increment `currentIndex` value independently.

```
public class Ledger {  
    public int currentIndex = 0 ;  
}
```

If we mark `currentIndex` as `Volatile` , then each thread will use its value from the main memory and will not create a copy of it.

```
public class Ledger {  
    public volatile int currentIndex = 0 ;  
}
```



Questions

- What is a volatile variable?
- Explain the problem that volatile field solves?
- Where are the volatile variables stored?

STATIC VS VOLATILE VS SYNCHRONIZED

static Variable

The **static** variables are used in the context of class objects where there is only one copy of a static variable exists, regardless of how many objects of the class are created.

But if multiple threads are accessing the same variable, each thread will make a copy of that variable in its CPU cache and change made by a thread will not be visible to other threads.

volatile Variable

volatile variables are used in the context of threads, where only one variable exists regardless of how many threads or objects access it and everyone always gets the most recently updated value. **Volatile** forces all the reads and writes to happen directly in the main memory and not in CPU caches.

synchronized

Both **static** and **volatile** are field modifiers dealing with memory visibility related to variables; whereas, **synchronized** deals with controlling access to a critical section in code using a monitor, thus preventing concurrent access to a section of code.

Questions

- What are the static variables?

- All the objects of a class share static variables. But in a multi-threaded environment; why a change made by one object to the static variable is not visible to an object on another thread?
- If both static and volatile variables are shared between objects, then what's the problem a volatile variable solves?
- How are volatile and static variables different from the synchronized, since even the synchronized monitor guards memory objects too?

THREADLOCAL STORAGE

Each thread has a private memory called **thread-local storage**, which can be used to store thread's current operation-related data. Usually, the **ThreadLocal** variables are implemented as private static fields and are used to store information like Transaction Identifier, User Identifiers, etc.

ThreadLocal declaration

```
ThreadLocal <String > threadLocal =  
    new ThreadLocal<String >();
```

Setting thread local value

```
threadLocal .set("Account id value" );
```

Getting thread-local value

```
String accountId = threadLocal .get();
```

- As ThreadLocal objects are contained within a thread, you don't have to worry about synchronizing the access to that object.
- ThreadLocal object existence is linked to the thread for which it was created, unless there are other variables which reference the same object too.
- To prevent leak, it's a good practice to remove ThreadLocal object using remove() method.

```
threadLocal .remove();
```

Questions

- What is thread-local storage?

- What are the things you should store in Thread Local Storage?
- Why shouldn't you synchronize the access to the objects that are stored in ThreadLocal Storage?
- Why should you remove object from Thread Local Storage?

WAIT() VS SLEEP()

wait()

Conditional Synchronization with wait().

```
public void manageWaitFor(int timeInMs)
    throws InterruptedException {
    synchronized (reference) {
        if (condition1) {
            // wait for notification
            reference.wait(timeInMs);
        }
    }
}
```

sleep()

Thread sleeping for a specified interval.

```
public void manageSleepFor(int timeInMs)
    throws InterruptedException {
    //Pause for timeInMs milliseconds
    Thread.sleep(timeInMs);
    //Print a message
    print("Slept for :" + timeInMs + "ms.");
}
```

- The **wait** method is called on the object's monitor; whereas, **sleep** is called on a **thread**.
- Waiting objects can be notified; whereas, a sleeping thread cannot.
- A sleeping thread cannot release a lock; whereas, awaiting object can.
- To wake a sleeping thread you need a reference of it, which is not needed for a waiting object.

Questions

- What is the difference between wait and sleep?
- Why it's possible to notify the waiting object but not the sleeping?
- As you need handle to a sleeping thread to wake it; do you need direct access to a waiting object too? Why?
- What is the mechanism to signal an object to come out of wait?

JOINING THREADS

Threads are usually joined when there is a dependency between the threads. The `join()` method of the target thread is used to suspend the current thread. In such situations a current thread cannot proceed, until the target thread on which it depends on, has finished execution.

```
// main thread joined with the thread
public void main(String[] args) {
    Thread thread1 = new Thread(
        ThreadMethodRef::
            threadMethod);
    thread1 .start();
    // current thread waits
    // until thread1 completes
    // execution
    thread1 .join();
}
```

Questions

- What is the purpose of the thread's join method?
- Why do you need to join two threads?

ATOMIC CLASSES

[Atomic classes](#) provide the ability to perform atomic operations on the primitive types, such that only one thread is allowed to change the value until the method call completes. Atomic classes like [AtomicInteger](#) and [AtomicLong](#) wraps the corresponding primitive types. There is one present for reference type too, [AtomicReference](#) .

There is no need to provide synchronized access to Atomic Class objects. Method [incrementAndGet\(\)](#) is [AtomicInteger](#) is often used in place of pre and post-increment operators.

Questions

- What are the atomic classes?
- Why don't you need to synchronize access to an object of the atomic class?
- Why pre and post-increment operator are not thread-safe?
 - *Pre and post-operation are multiple operations under the hood- read, increment, and write. All three are not synchronized together, so any thread context switch that happens in between will result in undesired results.*
- What is the difference between Atomic and Volatile variables?
 - *Atomic variables provide atomic access even for the compound operation like pre and post-increment operation, which is not possible if the variable is declared as Volatile. Volatile just guarantees happen-before reads.*

LOCK

- **Locking** is a mechanism for controlling shared resource access in a multithreaded system.
- **ReentrantLock** class implements a **lock interface** .
- A lock can be acquired and released in different blocks of code.
- **Lock** interface has method **tryLock()** to verify resource availability.
- As a good practice, the acquired lock must be released in the **finally** block.

```
// Thread-safe class
public class SafeAccount {
    // Create lock object
    private Lock lockObject =
        new ReentrantLock();
    public void addMoney() {
        // Acquire the lock
        lockObject .lock();
        try {
            // add some money logic here
        } finally {
            // Release the lock
            lockObject .unlock();
        }
    }
}
```

Questions

- Explain the locking mechanism in a multi-threaded environment?

- Do you need to acquire and release lock in the same block of code?
- Why should you prefer using `tryLock()` instead of `lock()`?

ReadWriteLock

- `ReadWriteLock` maintains a pair of associated locks, one for writing and the other for read-only operations.
- Only one thread can acquire a write lock, but multiple threads can have read lock.
- `ReadWriterLock` interface is implemented by `ReentrantReadWriteLock`.


```
// ReentrantReadWriteLock lock
ReentrantReadWriteLock rwl =
    new ReentrantReadWriteLock();
// read lock
Lock rl = rwl .readLock();
// write lock
Lock wl = rwl .writeLock();
```

Questions

- What are the benefits of using `ReadWriteLock`?
- In which scenarios should you prefer `ReadWriteLock` over other locking mechanism?

SYNCHRONISERS

Synchronizers synchronize multiple threads to protect a [Critical Section](#) .



Synchronizer Types

- Barriers
- Semaphore
- Phasers
- Exchangers
- Latch

Questions

- What is the purpose of synchronizers?
- What are the different types of synchronizers available in Java?

Barriers

- In **Barriers** , a set of threads waits for each other to arrive at the barrier point before moving ahead.
- **CyclicBarrier** is a concrete implementation of the **Barrier** synchronizer.

```
// barrier with five threads
CyclicBarrier barrier =
    new CyclicBarrier( 5 );
}
```

- A **Barrier** is also called a cycle because it can be reused after calling **reset()** on it.
- Action can be passed to the **CyclicBarrier** to execute when all the threads reach barrier point.

```
// barrier with an action to run
// at the barrier point.
CyclicBarrier barrier =
    new CyclicBarrier( 5 , () -> {
// barrier point action code.
});
```

- If any of the thread is terminated prematurely then all the other threads waiting at the barrier point will also exit.
- Barriers are generally used when you divide an operation into multiple tasks running on separate threads, and wait for all the tasks to complete before moving ahead.

Questions

- Explain Barrier synchronizer?
- Can you reuse the same Barrier object again? How?

- What if one of the threads, which other threads await a Barrier, dies?
- What job do you use Barrier for?

Semaphore

- Semaphore maintains a specified number of permits to access a Critical Section .

```
// Semaphore created with four permit  
// for four threads
```

```
Semaphore semaphore =  
    new Semaphore( 4 );
```

- To gain a permit, use the `acquire()` method. Each call to the `acquire()` method is blocked until the permit becomes available.
- To release the permit, use the `release()` method.
- A permit can be released by a different thread, other than the one that acquired it.
- If `release()` is called more times than `acquire()` , then for each such additional release, an additional permit will be added.
- If you wish to acquire a mutually exclusive lock, initialize the Semaphore with only one permit.
- Semaphore is generally used to allow limited access to an expensive resource.

Questions

- Explain Semaphore synchronizer?
- What happens if you call `release()` more time than `acquire()`?
- What happens when you call `acquire`, but permits are not available?
- How can you acquire a mutually exclusive lock using Semaphore?
- Where do you use Semaphore?

Phaser

- Unlike other barriers, the number of parties registered with the **Phaser** can dynamically change over time .

// Phaser with four registered parties

```
Phaser phaser =  
    new Phaser( 4 );
```

- A **phaser** can also be reused again.
- Use a **register()** method to register a party.
- When the final party for a given phase arrives, an optional action can be performed, and then the **Phaser** advances to the next phase.
- Use the **arriveAndAwaitAdvance()** method to wait for all parties to arrive before proceeding to the next phase.
- **Phasers** monitor the count of registered, arrived, and un-arrived parties. Even a caller who is not a registered party can monitor these counts on a **Phaser**.
- A party can be de-registered using the **arriveAndDeregister()** method, from moving to the next phase.

Questions

- Explain Phasor synchronizer?
- What is special about Phasor regarding the number of parties that can register with it?
- Can you reuse the same Phaser object again?
- Can you monitor the count of Phasor monitor using some external object who is not registered as a party with the Phaser?
- What is the difference between Semaphore and Phaser?

Exchanger

- The **exchanger** lets two threads wait for each other at a Synchronization point to swap elements.

```
// Exchanging array of strings
Exchanger<ArrayList<String>>
    exchanger =
    new Exchanger <ArrayList<String>>();
```

- Exchangers use **exchange()** method to exchange information.

```
// exchanger exchanging data
objectToExchange =
    exchanger .
    exchange( objectToExchange );
```

- On **exchange()** , the consumer empties the object to be exchanged and waits for the producer to exchange it with full object again.

Questions

- Explain Exchanger synchronizer?
- How many threads are required with the Exchanger object?
- What is the primary purpose of Exchanger synchronizer?
- Does Exchanger synchronizer use the same object to exchange every time or a different object can be exchanged?

Latch

- **Latch** makes the group of threads wait till a set of operations is finished.
- The latch cannot be reused.
- **CountDownLatch** class provides an implementation for **Latch** .
- All threads wait calling **await()** method till **countDown()** is called as many times the latch counter is set.

```
// Create a countdown latch with
// five counter
CountDownLatch cdl =
    new CountDownLatch( 5 );

// Count down on the latch after
// completion of thread job
cdl .countDown();

// waiting for count down signals
cdl .await();
```

Questions

- Explain Latch synchronizer?
- Can you reuse the same Latch object again; like Barrier and Phaser?
- What is the signaling process when a thread finishes its work?

EXECUTOR FRAMEWORK

The executor framework provides an infrastructure to execute a set of related tasks on a thread.

Executor takes care to manage the following.

- Creating and destroying threads.
- Maintaining an optimal number of threads for a task.
- Parallel and sequential execution of tasks.
- Segregating task submission and task execution.
- Policies are related to controlling task execution.

```
// Executor interface definition
public interface
    Executor {
    void execute ( Runnable command);
}
```

Questions

- Explain the Executor framework?
- What are the different capabilities of the Executor framework?

Executor Service

- ExecutorService inherits from the Executor interface providing the following additional methods.
 - `shutdown()` - shuts down the executor after submitting the tasks.
 - `shutdownNow()` - interrupts the current task and discards the pending tasks.
 - `submit()` - adds tasks to the `Executor` .
 - `awaitTermination()` - waits for existing tasks to terminate.
- ExecutorService provides Future objects to track the progress and the status of the executing task.
- All the tasks submitted to the Executor are queued, which are executed by the thread pool threads.

// Executor created with five threads in its thread pool

```
ExecutorService exec =  
    Executors.  
        newFixedThreadPool( 5 );
```

- To create a thread pool with a single thread, use `newSingleThreadExecutor()` method.

Handling Results

- The `run()` method of the `Runnable` interface cannot return a result or throw an exception.
- Tasks, which can return a result, are instance of a `Callable` interface.

//tasks can return results derived from Callable using the call method

```
public interface Callable < V > {  
    V call() throws Exception;  
}
```

- The `submit()` method returns `Future` object which helps to track task.

```
// ExecutorService example
public class ExecService {
    public static void main(String[] args)
        throws ExecutionException,
        InterruptedException {
        // Create executor with five threads
        // in its thread pool.
        ExecutorService exec =
            Executors.newFixedThreadPool( 5 );
        // Submit the callable task to the executor
        Future <String> task =
            exec .submit(
                new Callable <String>() {
                    @Override
                    public String call()
                        throws Exception {
                        //some logic
                        return null ;
                    }
                });
        // waits for the result
        String result =
            task .get();
        // Shutdown executor
        exec .shutdown();
    }
}
```

- If there is an exception during the task execution, calling `get()` method on the `ExecutorService` will throw an instance of `ExecutionException`.

Scheduling Task

- `ScheduledExecutorService` can be used to schedule a future task.
- **M**ethods used to schedule task.
 - `schedule(`
 `task ,`
 `delayTime ,`
 `TimeUnit.SECONDS)`
 - `scheduleAtFixedRate(`
 `task ,`
 `delayTime ,`
 `repeatPeriod ,`
 `TimeUnit.SECONDS)`
 - `scheduleWithFixedDelay(`
 `task ,`
 `delayTime ,`
 `fixedDelay ,`
 `TimeUnit.SECONDS);`

ExecutorCompletionService

- `ExecutorCompletionService` uses `Executor` to execute the task.
- `CompletionService` of `Executor` can be used to get results from multiple tasks.
- `ExecutorCompletionService` provides concrete implementation for `CompletionService`.

```
// Create executor with five threads
ExecutorService es =
    Executors.newFixedThreadPool( 5 );

// ExecutorCompletionService returns an object
ExecutorCompletionService<Result> cs =
    new ExecutorCompletionService<>( es );

// submit task to ExecutorCompletionService
cs .submit( longTask );

// get the result of task
Future <Result> completedTask =
```

```
cs .take();
```

Notes

- Always associate context-based names to the threads, it immensely helps in debugging.
- Always exit gracefully, by calling either `shutdown()` or `shutdownNow()` based on your use case.
- Configure thread pool for the `ExecutorService` such that the number of threads configured in the pool are not significantly greater than the number of processors available in the system.
- You should query the host to find the number of processors to configure the thread pool.

```
Runtime.getRuntime().  
    availableProcessors();
```

Questions

- Explain `ExecutorService`?
- How do you track the progress and status of the executing tasks?
- Does `Executor` service use dedicated threads to execute queued tasks?
- Can you use `Runnable` object with `ExecutorService`? Why?
- How do you find if an exception is thrown in the `ExecutorService`?
- Can you schedule a task to run in the future with the `ExecutorService`?
- With `ExecutorService`, how can you get results from the multiple tasks?
- What is the difference between `submit()` and `execute()` methods of `ExecutorService`?

- *If you use submit(), you can get any exception thrown by calling get() method on Future; whereas, if you use execute(), an exception will go to `UncaughtExceptionHandler` .*

- How can you exit gracefully from ExecutorService?
- What should be the criteria for configuring thread pool size? How can you set that?

FORK-JOIN

- The `fork-join` framework takes advantage of multi-processors and multi-core systems.
- It divides the tasks into sub-tasks to execute in parallel.
- The `fork()` method spawns a new sub-task from the task.

```
// spawn subtask  
subTask .fork();
```

- The `join()` method lets the task wait for other task to complete.

```
// wait for subtask to complete  
subTask .join();
```

- Important classes in Fork-Join
 - `ForkJoinPool` - thread pool class is used to execute subtasks.
 - `ForkJoinTask` - manages subtask using `fork()` and `join()` methods.
 - `RecursiveTask` - a task that yields result.
 - `RecursiveAction` - a task that does not yield results.
- Both `RecursiveTask` and `RecursiveAction` provide an abstract `compute ()` method to be implemented by the class, whose object represents the `ForkJoin` task.

Questions

- Explain the Fork-Join framework in Java?
- How does the Fork-Join framework help to optimize task execution?

- What is the difference between RecursiveTask and RecursiveAction?

REFLECTION

REFLECTION

[Reflection](#) is used to examine the code runtime and possibly modify the runtime behavior of an application.

Purpose of reflection

[Reflection](#) must be used only for special-purpose problem solving and only when information is not publicly available. The class members are marked private for reasons. Few of the popular usage of reflection in day to day development includes the following:

- Reflection enables modular architecture by investigating code and libraries runtime to plugin classes and components in the application.
- Annotations are read using reflection. [JUnit](#) uses reflection to discover methods to set up and test.
- Debuggers use reflection to read private members of the class.
- IDEs use reflection to enumerate class members and probe code.
- Object-relational mappers use reflection to create objects from data.
- Dependency Injection framework like [Spring](#) uses reflection to resolve dependencies.

Usage

The code below demonstrates the use of reflection to access the private field of class [Account](#) .

```
public class Account {  
    private float rate = 10.5f ;  
}
```

```

public static void main( String [] args)
    throws IllegalAccessException {

    Account account = new Account();
    Class <? extends Account > refClass =
        account .getClass();

    // get all the field
    Field [] fields =
        refClass .getDeclaredFields();

    for ( Field field : fields ){
        // no more private
        field .setAccessible( true );
        print( field .get( account ));
    }
}

```

Questions

- What is reflection?
- What are the different scenarios where reflection can be used?
- How do debuggers use reflection?
- How can reflection help in creating modular software architecture?
- Can you access private members using reflection?
- Give some examples, where reflection is used in software development libraries, tools, and frameworks?

DRAWBACKS OF REFLECTION

[Reflection](#) must be used only when something is not publicly available and there is a very compelling reason, otherwise, you must review your design.

Drawbacks

- [Reflection](#) does not have compile-time checking, any change in member name will break the code.
- [Reflection](#) violates encapsulation as it reveals the internal data structures.
- [Reflection](#) violates abstraction as it reveals the internal implementation and provides the ability to bypass validations applied to the members.
- [Reflection](#) is slow, as it has additional overhead at runtime to resolve the members.

Questions

- When should you use reflection?
- How does reflection violate encapsulation?
- How does reflection violate abstraction?
- Why reflection is considered slow?
- Why reflection code is considered fragile and can break?

DATA INTERCHANGE

JSON

JavaScript Object Notation or JSON has become extremely popular for the data interchange in the last few years. Now it's not just an alternative to XML but is a successor to it. With the evolution of Big Data and the Internet of Things, along with the JSON's ability to be easily parsed to JavaScript objects, it has become a preferred data format for integration with the web. The following are the few primary factors behind this.


- It's interoperable, as it's restricted to primitive data types only.
- It's lightweight and less verbose than XML.
- It's very easy to serialize and transmit structured data over the network.
- Almost all modern languages support it.
- JavaScript parser in all the popular web browsers supports it.

JSON Structures

JSON structures can be categorized as a JSON Object and JSON array.

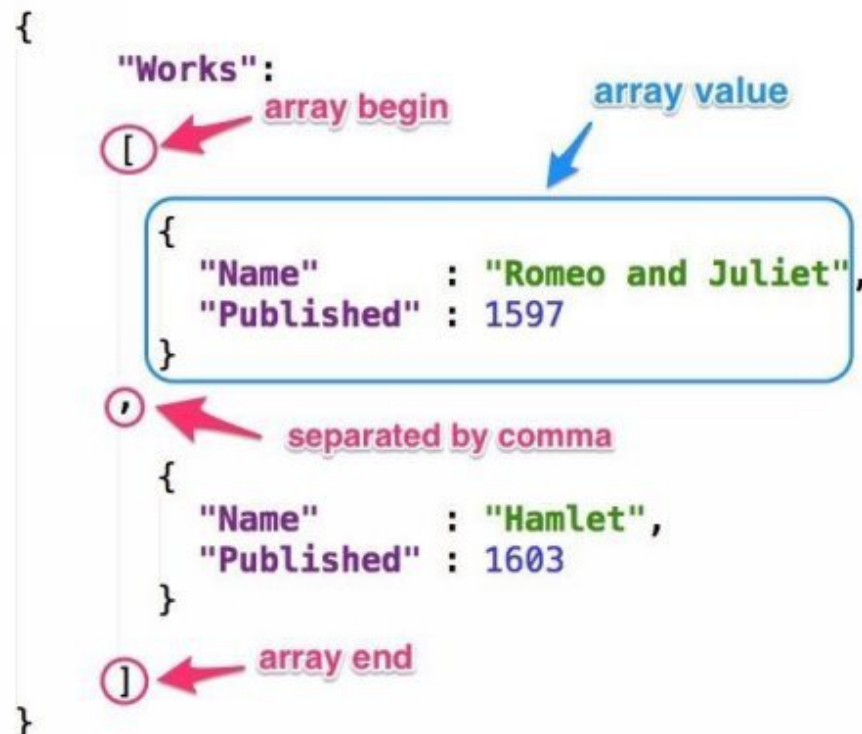
JSON Object

A JSON object is simply a name-value pair separated by a comma. Name is always a string.



JSON Array

JSON Array is an ordered collection of values. Value can be a string, a number, an object, or an array itself.



Questions

- What is JSON?
- What is the main reason for JSON's popularity?
- What are the significant differences between JSON and XML data formats?
- Why JSON has become the preferred format for data interchange?
- Explain the difference between JSON Object and JSON Array?

MEMORY MANAGEMENT


STACK VS HEAP

- A [stack](#) is a memory associated with each system thread when it's created; whereas, a heap is shared by all threads in an application.
- For each function, the thread visits, a block of memory is allotted on the top of the stack, for local variables and bookkeeping data, which gets freed when that function returns, in a LIFO order. In contrast, the allocation of memory in Heap is relatively random with no enforced pattern, and variables on the heap are destroyed manually.
- When the thread exits, the stack associated with the thread is reclaimed. When the application process exists, heap memory is reclaimed.
- Allocating and freeing stack memory is simpler and quicker as it is as simple as adjusting pointer. Allocating and freeing memory is comparatively complex in Heap, as there is no fixed pattern of memory allocation.
- The stack memory is visible only to the owner thread, so memory access is straight forward; whereas, heap memory is shared across multiple threads in the application, so synchronization with other thread has performance implications.
- When stack memory is exhausted, JVM throw [StackOverflowError](#) ; whereas, when heap space is exhausted, JVM throws [OutOfMemoryError](#) .

Memory allocation in stack and heap

In this example, stack and heap memory allocation is depicted, when a method is invoked.

- [num1](#) and [num2](#) variables are stored in a stack.
- The reference [point](#) variable is stored in a stack.
- An object of class [Point](#) is stored in heap.



Questions

- What is a stack?
- What is a heap?
- During its lifecycle how is the memory allocated and de-allocated in a stack?
- When does the stack memory gets released?
- When does the heap memory gets released?
- Why memory allocation in a stack is faster compared to the heap?
- Why do you need to apply synchronization to the heap memory and not to the stack memory?
- What exception do you get when stack memory is exhausted?
- What exception do you get when heap memory is exhausted?
- Do you need to explicitly release the objects on the stack?

HEAP FRAGMENTATION

[Heap fragmentation](#) happens when a Java application allocates and de-allocates small and large blocks of memory over a period, which leads to lots of small free blocks of memory spread between used blocks of memory. This may lead to a situation when there is no space left to allocate a large block of memory, even though the cumulative size of entire small free blocks is more than the required memory for the large block.

[Heap fragmentation](#) causes a long [Garbage Collection](#) cycle as JVM is forced to compact the heap. Avoiding allocating a large block of memory, by increasing heap size, etc, can control [heap fragmentation](#) .

Questions

- What is heap fragmentation?
- Why does heap fragmentation happen?
- Who has the responsibility to reduce heap fragmentation?
- How can you control heap fragmentation?
- What happens when there is a very high level of heap fragmentation?
- Why does heap fragmentation slow down the application?

OBJECT SERIALIZATION

Converting the content of an in-memory object into bytes to either persist it or to transfer, is called [object serialization](#) . These bytes can be converted back to object by de-serializing it.

In Java, an object is serializable if its class implements [java.io.Serializable](#) or its subinterface [java.io.Externalizable](#) . Members marked as [transient](#) are not serialized. [ObjectInputStream](#) and [ObjectOutputStream](#) are stream classes specifically used to read and write objects.

Questions

- What are serialization and de-serialization?
- Which interface needs to be implemented by the type that want's to support serialization?
- How can you prevent a member from serialization?
- Which classes in Java are used to serialize and de-serialize objects?

GARBAGE COLLECTION

Garbage Collection in Java is the process to identify and remove the unreferenced objects from the memory and also move the remaining objects together to release a contiguous block of memory. When Garbage collection happens, all the running threads in the application are suspended during the collection cycle. Garbage Collector runs on a Daemon thread.

Moving all the surviving objects together reduces the memory fragmentation, thus improving the performance of memory allocation.

During the Garbage Collection cycle, objects are moved to different areas in memory, known as **generations** , based on their survival age.

System class exposes method **gc()**, which can be used to request Garbage Collection. When you call **System.gc()** , JVM does not guarantee to execute garbage collection immediately but may perform when it can. You can also use **Runtime.getRuntime().gc()** to request Garbage Collection.

Questions

- What is Garbage Collection?
- Explain the Garbage Collection cycle?
- Why Garbage Collection is considered an expensive process?
- Why Garbage Collection cycle improves the performance of an application?
- How does Garbage Collection prevent OutOfMemoryException?
- Which type of thread is used for Garbage Collection: Daemon or user thread?

- What is memory fragmentation?
- Why does the application seem to slow down when Garbage Collection happens?
- Can you explicitly request a Garbage Collection?
- Which methods can you call to request Garbage Collection?
- Does Garbage Collection always happen when requested?

MEMORY MANAGEMENT

JVM memory is divided into two major categories, [stack memory](#), and [heap memory](#).


[Stack Memory](#)

[Stack memory](#) is associated with each system thread and used during the execution of the thread. [A stack](#) contains local objects and the reference variables defined in the method; although the referenced objects are stored in heap. Once the execution leaves the method, all the local variables declared within the method are removed from the stack.

[Heap Memory](#)

[Heap Memory](#) is divided into various regions called generations:

- [New Generation](#) - It's divided into Eden and Survivor space. Most of the new objects are created in Eden's memory space.
- [Old Generation](#)
- [Metaspace](#)



When the **new generation** is filled, it triggers garbage collection. Objects that survive this GC cycle in **Eden** are moved to **Survivor** space. Similarly, after a few cycles of GC, the surviving objects keep moving to the **old generation**. **Metaspace** is used by JVM to keep permanent objects, mostly the metadata information of the classes. New generations are more frequently garbage collected than the old generations.

- **New Gen** and **Old Gen** are part of heap whereas **Metaspace** is part of Native memory.
- **Metaspace** can expand at runtime, as it's part of native memory.

When Garbage collection happens, all the application threads are frozen until GC completes its operation. Garbage collection is typically slow in the old generation; so if lots of Garbage Collection happens in the old generation, it may lead to timeout errors in the application.

Questions

- How is Heap memory divided?
- Explain different generations of heap memory? How objects are moved across generations?
- How is the garbage collection cycle triggered?

- What are Eden and Survivor spaces in New Generation memory?
- When do the objects in Eden move to Survivor space?
- What is Metaspace?
- What is the type of things stored in Metaspace?
- Why Metaspace has virtually unlimited space?
- In which generation Garbage Collection cycle runs slowest?
Why?
- Which generation is more frequently garbage collected?

WEAK VS SOFT VS PHANTOM REFERENCE

Generally, if an object is having a reference then the garbage collector will not collect it. This principle is not true for Weak and Soft references.

Weak Reference

A **weak reference** is a reference that eagerly gets collected by the garbage collector. Weak references are good for caching, which can be reloaded when required.

Soft Reference

Soft reference is slightly stronger than the Weak Reference, as these are collected by garbage collector only when there is a memory constraint. Soft references are generally used for caching re-creatable resources like file handles, etc.

Phantom Reference

Phantom reference is the weakest reference in Java. It is referenced after an object has been finalized, but the memory is yet to be claimed by the Garbage collector. It is primarily used for the technical purpose to track memory usage.

Questions

- What is a weak reference?
- What is a soft reference?
- What is Phantom Reference?

- What are the various reference types that get collected by the Garbage Collector, even when the objects of their types are still in use?
- What is the difference between soft reference and weak reference?
- What is the typical usage of weak reference?
- What is the typical usage of soft reference?
- What is the typical usage of Phantom reference?
- Which is collected first: soft reference or weak reference?

UNIT TESTING

WHY UNIT TESTING?

Two most important reasons

- You can understand the real benefits only by doing it yourself.
- The unit test helps you to sleep well at night.

Other important reasons

- Unit tests provide immediate and continuous feedback on the success/failure of the changes made to the code.
- Units test increases the confidence to make big changes without worrying about breaking any existing feature.
- The unit test helps you to understand the internals of the code and design.
- Unit tests also serve as documentation for various coding scenarios.
- Unit test saves time in the longer run by reducing the multiple cycles of manual verification of different scenarios.

Questions

- What are the benefits of unit testing?
- How does unit testing save time in the long run?
- How can you use unit testing to document the coding scenarios?

UNIT VS INTEGRATION VS REGRESSION VS VALIDATION TESTING

Unit

Unit testing is continuously done while writing the code; to get immediate feedback to the smallest testable change made. Smallest testable units can span across methods or classes but must exclude external dependencies like File I/O, Databases, Network Access, etc.

Integration

Integration testing is done to test end to end integration when all the changes made for a scenario is completely implemented.

Regression

Regression testing is a series of tests performed on entire software to uncover bugs in both functional and non-functional areas. Regression testing is usually done after enhancements, software updates, etc.

Validation

Validation testing is done to verify that after updating or deploying the software, the changes are made as per the requirements.

Questions

- What is unit testing?

- What is integration testing?
- What is regression testing?
- How are the differences between integration testing and regression testing?
- What is validation testing?

TESTING PRIVATE MEMBERS

Private members can be tested using [reflection](#), but it's advisable to do so only when you need to test some legacy code, where changing the visibility of the private method is not allowed.

Usage

The code below demonstrates the use of reflection to access the private field of class [Account](#) .

```
public class Account {  
    private float rate = 10.5f ;  
}  
  
@Test  
public void testConcatenate() throws  
    IllegalAccessException ,  
    NoSuchFieldException {  
  
    Account account = new Account();  
    Class <? extends Account > refClass =  
        account .getClass();  
    Field field =  
        refClass .getDeclaredField( "rate" );  
    field .setAccessible( true );  
  
    assertEquals( 10.5f , field .get( account ));  
}
```

As mentioned, it's not advisable to access private fields using reflection; we have following other options to test code in a private method:

- Test the private method through the public method.
- Change the access modifier of the field, if possible.
- Change the class design.

Questions

- What are the different ways to test private members?
- Can you use Reflection to test private members?
- Why shouldn't you use reflection to test private members?
- What are the various methods to test private members?

MOCKING

The primary responsibility of a unit test is to test the conditional logic in the class code, which should run super-fast for immediate feedback. To enable immediate feedback, a class must have no external dependencies, which is not practical in object-oriented software development, where you need to communicate with external objects to perform File IO, read and write to a database, get data from web service, etc. So the basic idea of **mocking** is to replace these external dependencies with mock objects, to isolate the object under test.

Benefits of mocking

- Mock objects help you to isolate and test only the conditional logic in the class without testing dependencies.
- In the mock object, you can implement partial functionality required for the test without implementing the complete dependency object.
- You don't have to worry about understanding the internal of dependencies, which helps in faster development time.

Questions

- What are mock objects?
- What are the benefits of using mock objects?
- Why do mock objects help to speed up the tests?
- Why should you use mock objects for unit testing?
- What are the types of dependencies that you should replace with mock objects?

JAVA DEVELOPMENT TOOLS

GIT

Git is a source code management system. Its **distributed version control system (DVCS)** facilitates multiple engineers and teams to work in parallel. Git's primary emphasis is on providing speed with maintaining data integrity.

Git also provides the ability to perform almost all the operations offline, when network is not available. All the changes can be pushed to the server on the network's availability.

Let's discuss a few terms that are frequently used:

Repository - is a directory that contains all the project files.

Clone - it creates a working copy of the local repository.

Branch - is created to encapsulate the development of a new feature or to fix a bug.


HEAD - points to the last commit.

Commit - commits changes to the HEAD and not to the remote repository.

Pull - Gets the changes from the remote repository to the local repository.

Push - Commits the local repository changes to the remote repository.

Git Workflow Structure



Questions

- What is distributed version control systems?
- Why do you use a version control system?
- What are the typical activities you perform with a version control system?
- Explain the branching strategy of code?

MAVEN

Maven is a software project management framework that manages project build, dependency resolution, testing, deployment, reporting, etc. Maven is based on conventions rather than elaborate configurations.

Maven is managed by **pom.xml** (also known as Project Object Model) file, which facilitates defining various tasks (or goals) and also has a dedicated section to specify various project dependencies, which are resolved by Maven.

Questions

- Why should you need a framework for dependency resolution - you can manually download the files and add a reference?
- What all things does Maven support?
- Why Maven is said to be convention-based?
- What is the meaning of goals in Maven?

JENKINS

Continuous Integration

Continuous Integration is a practice where the engineers, who are working in parallel on the same code repository, merge their changes frequently. Every check-in is followed by a series of automated activities, which aims to validate the changes in the check-in. Typical automated activities following the check-in are:

- Fetching changes from the repository,
- Performing automated build.
- Execute different tests like the unit, integration, validations, etc.
- Deploy the changes.
- Publish the results.

Jenkins

Jenkins is an open-source web-based tool to perform continuous integration. Jenkins provides configuration options to configure and execute all the above-mentioned activities. Available configuration options are: configuring JDK, security, Build Script; integration with Git, Ants, Maven, Gradle, etc.; deployment, etc.

The execution of jobs can be associated with some event or can be based on time-based scheduling.

Questions

- What is Continuous Integration?
- What are the different activities that the Continuous Integration framework should support?

- Explain the capabilities of Jenkins to support Continuous Integration?
- How do you integrate third-party tools with Jenkins?

JAVA INTERVIEW COURSE